

CS425

Computer Systems Architecture

Fall 2025

**Re-Order Buffer:
Precise Exceptions and Speculation**

Exception Behavior with ROB

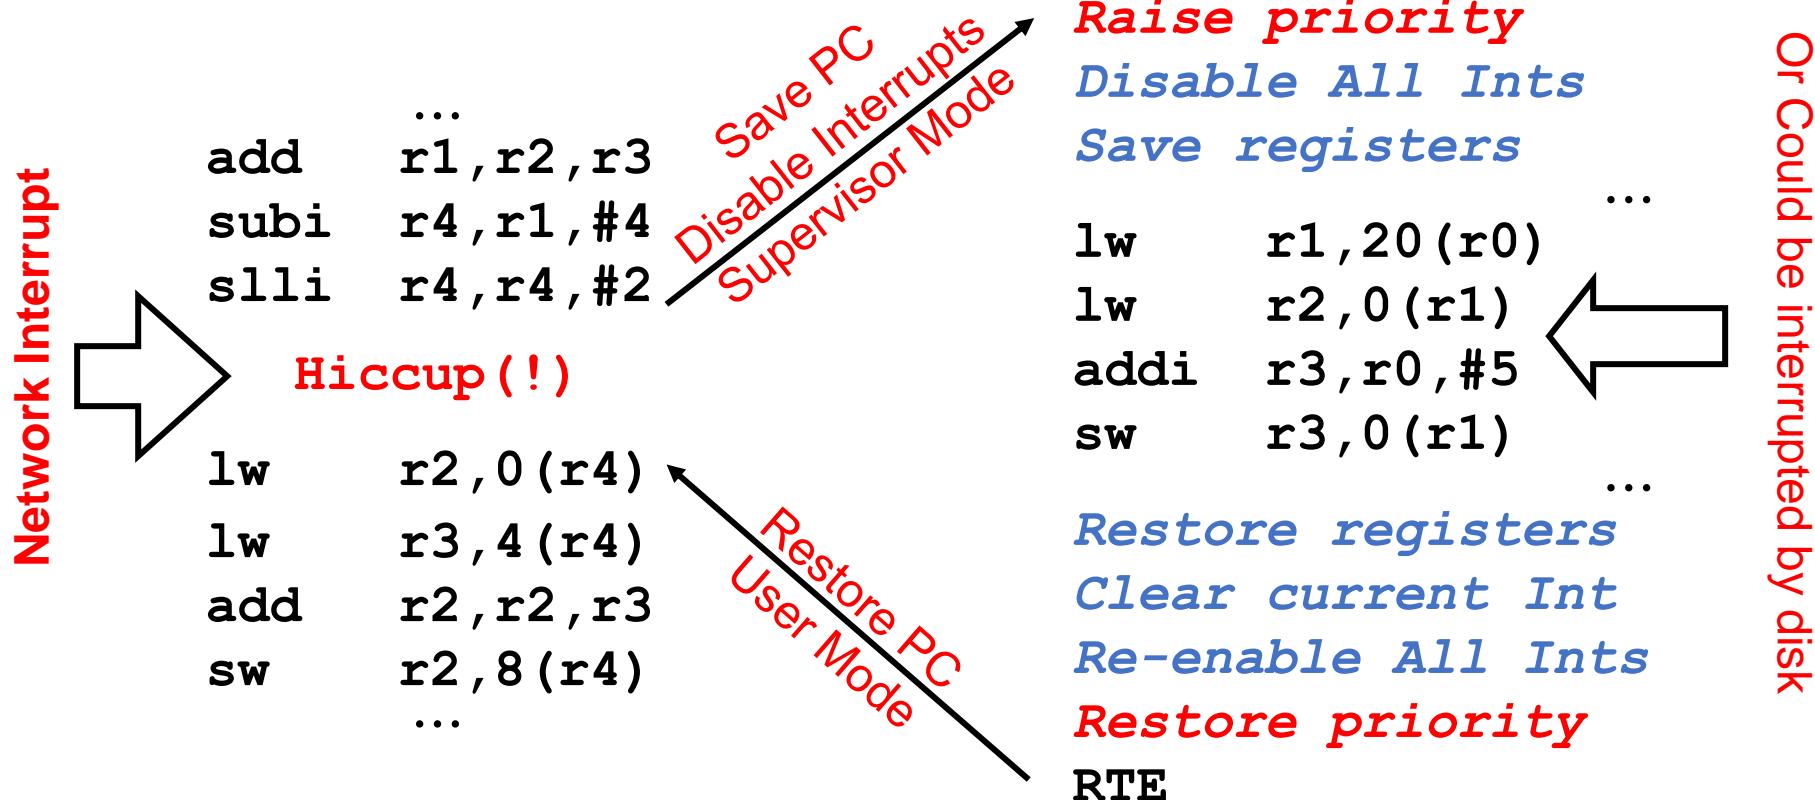
$$\text{CPI} = \text{CPI}_{\text{IDEAL}} + \text{Stalls}_{\text{STRUC}} + \text{Stalls}_{\text{RAW}} + \text{Stalls}_{\text{WAR}} + \text{Stalls}_{\text{WAW}} + \text{Stalls}_{\text{CONTROL}}$$

- Have to maintain:

- Data Flow
- Exception Behavior

Dynamic instruction scheduling (HW)	Static instruction scheduling (SW/compiler)
Scoreboard (reduce RAW stalls)	Loop Unrolling
Register Renaming (reduce WAR & WAW stalls) • Tomasulo • Reorder buffer	SW pipelining
Branch Prediction (reduce control stalls)	Trace Scheduling
Multiple Issue (CPI < 1)	
Multithreading (CPI < 1)	

Device Interrupt



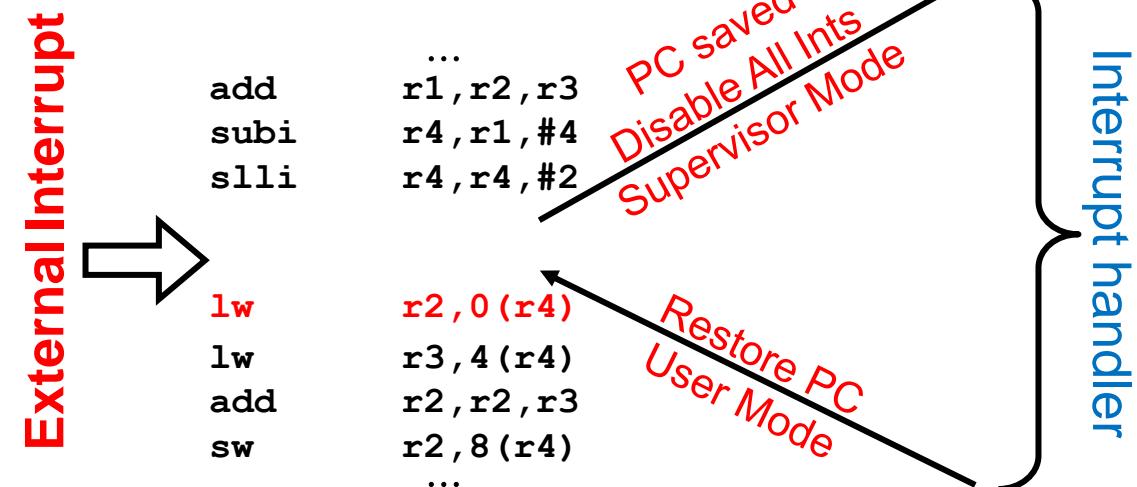
Note that priority must be raised to avoid recursive interrupts!

Types of Interrupts/Exceptions

- I/O device request
- Invoking an operating system service from a user program
- Breakpoint (programmer-requested interrupt)
- Integer arithmetic overflow
- FP arithmetic anomaly
- Page fault (not in main memory)
- Misaligned memory accesses (if alignment is required)
- Memory protection violation
- Using an undefined or unimplemented instruction
- Hardware malfunctions
- Power failure

Precise Interrupts/Exceptions

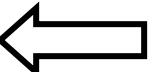
- An interrupt or exception is precise if there is an instruction (or interrupt point) for which:
 - All instructions before this instruction have fully completed
 - None of the instructions after this instruction (including the interrupting instruction) has modified the machine state
- This means that we can restart the execution from the interrupt point and still “get the correct results”
 - In the example: the Interrupt point is the `lw` instruction



Imprecise Interrupt/Exception

- An exception is imprecise if the processor state when an exception is raised does not look **exactly** as if the instructions were executed sequentially in strict program order
- Occurrence in two possibilities:
 - The pipeline may have already completed instructions that are later in program order
 - The pipeline may have not yet completed some instructions that are earlier in program order

Precise interrupt point requires multiple PCs when there are delayed branches

<pre>addi r4,r3,#4 sub r1,r2,r3 PC: bne r1,there PC+4: and r2,r3,r5 <other insts></pre>	 Interrupt point described as <PC, PC+4>
<pre>addi r4,r3,#4 sub r1,r2,r3 PC: bne r1,there PC+4: and r2,r3,r5 <other insts></pre>	 Interrupt point described as: <PC+4, there> (branch was taken) or <PC+4, PC+8> (branch was not taken)

Why do we need precise interrupts?

- Several interrupts/exceptions need to be restartable
 - i.e. TLB faults. Fix translation and then restart the faulting load/store
 - IEEE gradual underflow, illegal operation,

e.g: $f(x) = \frac{\sin(x)}{x}$

$x \rightarrow 0 \quad f(0) = \frac{0}{0} \Rightarrow NaN + illegal_operation$

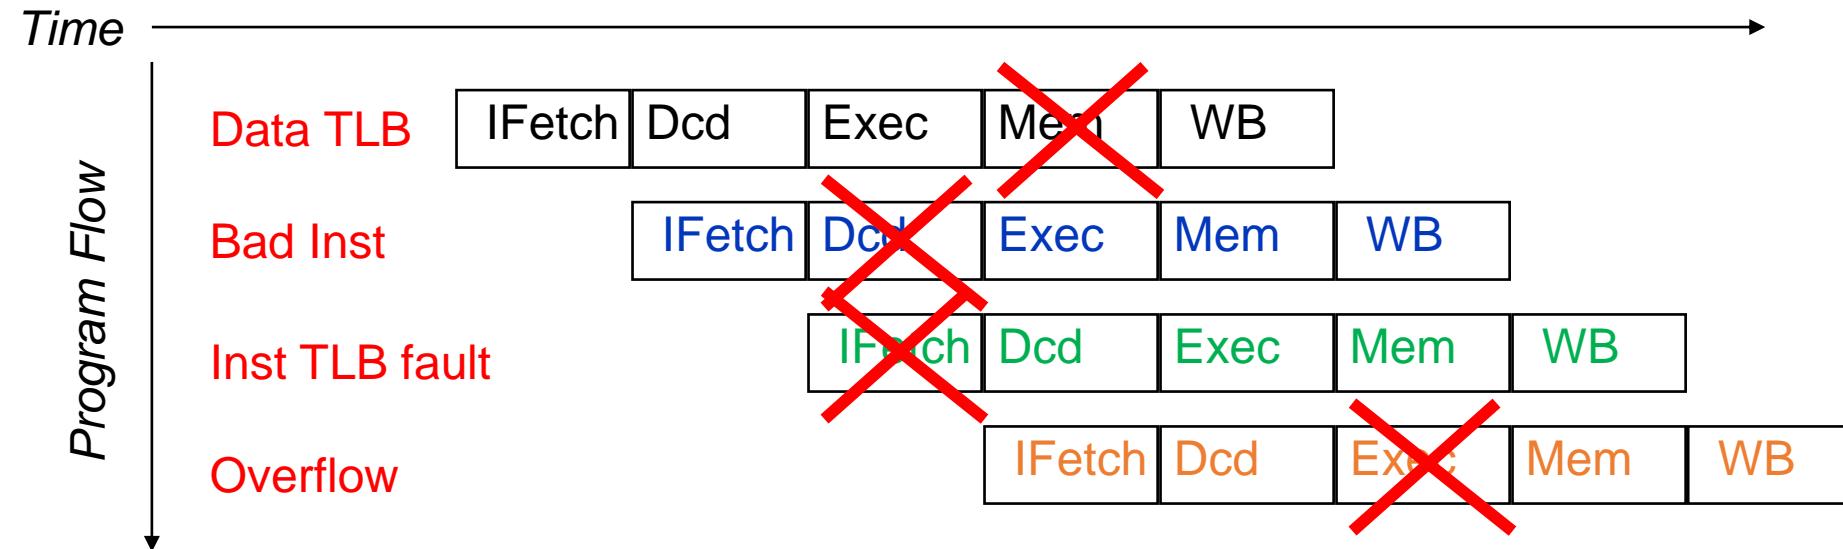
Want to take exception, replace *Nan* with 1, then restart.

- Restartability does not require *preciseness*. However, *preciseness* makes restarts *much simpler*
- Simplifies the Operating System (OS)
 - *Less state* needs to be saved away if unloading process.
 - Quick to restart (for fast interrupts)

Precise Exceptions in 5-stage RISC

- Exceptions may occur in different stages of the processor pipeline (i.e. out of order):
 - Arithmetic exceptions occur in execution stage
 - TLB faults can occur in instruction fetch or memory stage
- How do we guarantee precise exceptions? Mark the instructions with an “exception status” and wait until the WB stage to take the exception
 - Interrupts are marked as NOPs (like bubbles) that are placed into pipeline instead of an instruction.
 - Assume that interrupt condition persists in case NOP flushed
 - Clever instruction fetch might start fetching instructions from interrupt vector, but this is complicated and needs to switch to supervisor mode, saving of one or more PCs, etc

Another look at the exception problem



- Use the pipeline!
 - Each instruction has an exception status field.
 - Keep the PCs for every instruction in the pipeline.
 - Check the exception status when the instruction reaches the WB stage
- When an instruction reaches the WB stage and has an exception:
 - Save PC \Rightarrow EPC, Interrupt vector addr \Rightarrow PC
 - Convert all fetched instructions to NOPs
- It works because of in-order completion/WB

Scoreboard Example: Cycle 62

Instruction status:

Instruction	<i>j</i>	<i>k</i>	Issue	Read	Exec	Write
				<i>Oper</i>	<i>Comp</i>	<i>Result</i>
LD	F6	34+	R2	1	2	3
LD	F2	45+	R3	5	6	7
MULTD	F0	F2	F4	6	9	19
SUBD	F8	F6	F2	7	9	11
DIVD	F10	F0	F6	8	21	61
ADDD	F6	F8	F2	13	14	16

In-order issue

Out-of-order execute

Out-of-order commit!

Functional unit status:

Time	Name	dest		<i>S1</i>	<i>S2</i>	<i>FU</i>	<i>FU</i>	<i>Fj?</i>	<i>Fk?</i>
		<i>Busy</i>	<i>Op</i>	<i>Fi</i>	<i>Fj</i>	<i>Fk</i>	<i>Qj</i>	<i>Qk</i>	<i>Rj</i>
	Integer	No							
	Mult1	No							
	Mult2	No							
	Add	No							
	Divide	No							

Register result status:

Clock	<i>F0</i>	<i>F2</i>	<i>F4</i>	<i>F6</i>	<i>F8</i>	<i>F10</i>	<i>F12</i>	...	<i>F30</i>
62	<i>FU</i>								

Tomasulo Example: Cycle 57

Instruction status:

Instruction	<i>j</i>	<i>k</i>	Issue	Exec	Write	Busy	Address
				Comp	Result		
LD	F6	34+	R2	1	4	Load1	No
LD	F2	45+	R3	2	5	Load2	No
MULTD	F0	F2	F4	3	16	Load3	No
SUBD	F8	F6	F2	4	8		
DIVD	F10	F0	F6	5	57		
ADDD	F6	F8	F2	6	11		

Reservation Stations:

Time	Name	Busy	Op	S1	S2	RS	RS
				Vj	Vk	Qj	Qk
	Add1	No					
	Add2	No					
	Add3	No					
	Mult1	No					
	Mult2	Yes	DIVD	M*F4	M(A1)		

Register result status:

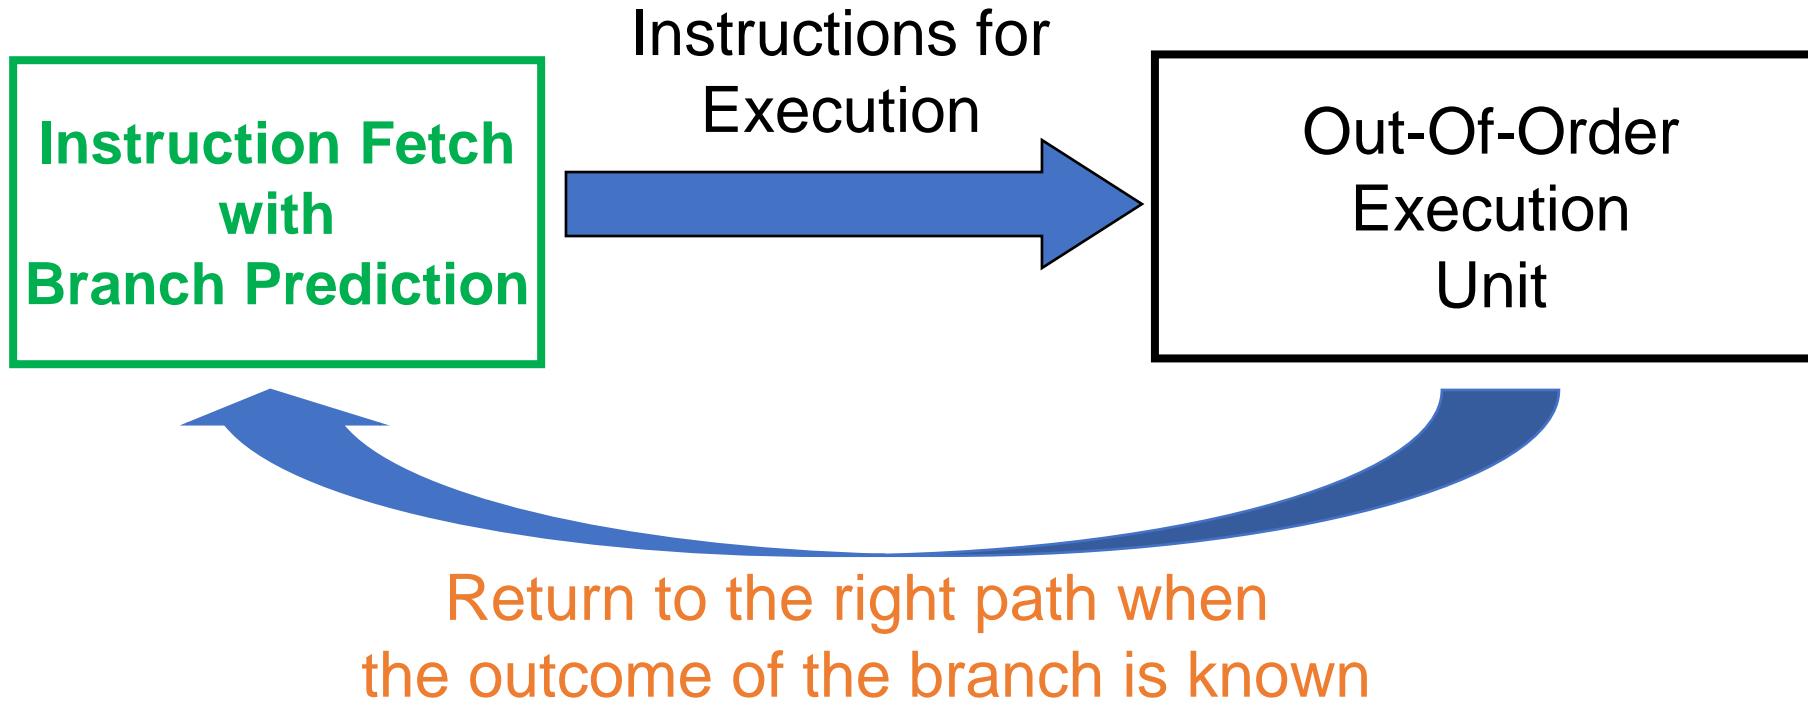
Clock	F0	F2	F4	F6	F8	F10	F12	...	F30
56	FU	Result							
	Reg File M*F4 M(A2) (M-M+M) (M-M)								

In-order issue

Out-of-order execute

Out-of-order commit!

Issue: “Fetch” unit



- Instructions from a potentially mispredicted branch path have been already executed.
- Instruction fetch decoupled from execution

Branch must execute fast for loop overlap!

- In the loop-unrolling example, we assume that the branches are executed from a “fast” integer unit to achieve overlap!

Loop :	LD	F0	0	R1
	MULTD	F4	F0	F2
	SD	F4	0	R1
	SUBI	R1	R1	#8
	BNEZ	R1	Loop	

- What happens if the branch depends on the outcome of MULTD?
 - We lose all benefits
 - We have to predict the outcome of the branch
 - If we predict “taken” the prediction would be correct most of the time.

Prediction: Branches, Dependencies, Data

- Branch Prediction is necessary for good performance
- We studied branches in the previous lecture. Modern architectures now predict many things: **data dependencies**, **actual data**, and **results of groups of instructions**
- Why does prediction work?
 - Underlying algorithm has regularities.
 - Data that is being operated on has regularities.
 - Instruction sequence has redundancies that are artifacts of way that humans/compilers think about problems.

Problem: Out-of-Order Completion

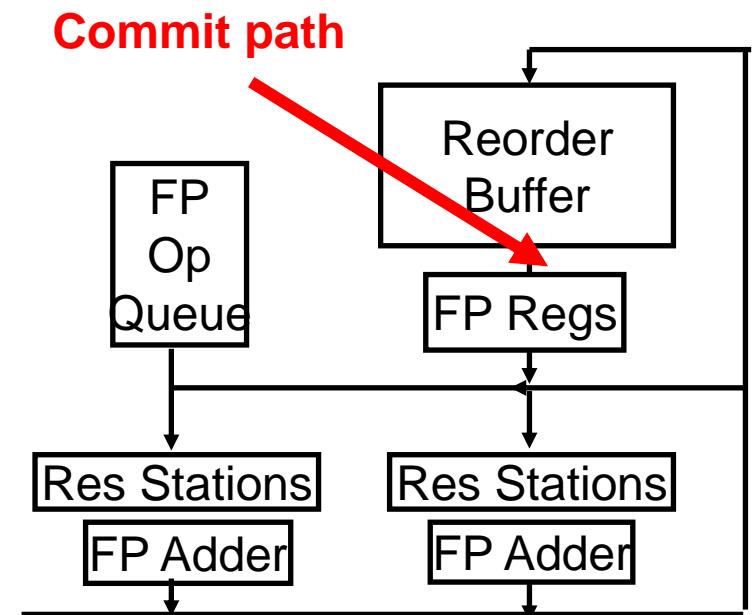
- Scoreboard and Tomasulo operate as follows:
 - In-order issue, out-of-order execution, **out-of-order completion**
- We need a way to synchronize the completion stage of instructions with the program order (i.e. with issue-order)
 - Easiest way is with **in-order completion** (i.e. re-order buffer)
 - Other Techniques (Smith paper): Future File, History Buffer

Precise Interrupts and Speculation:

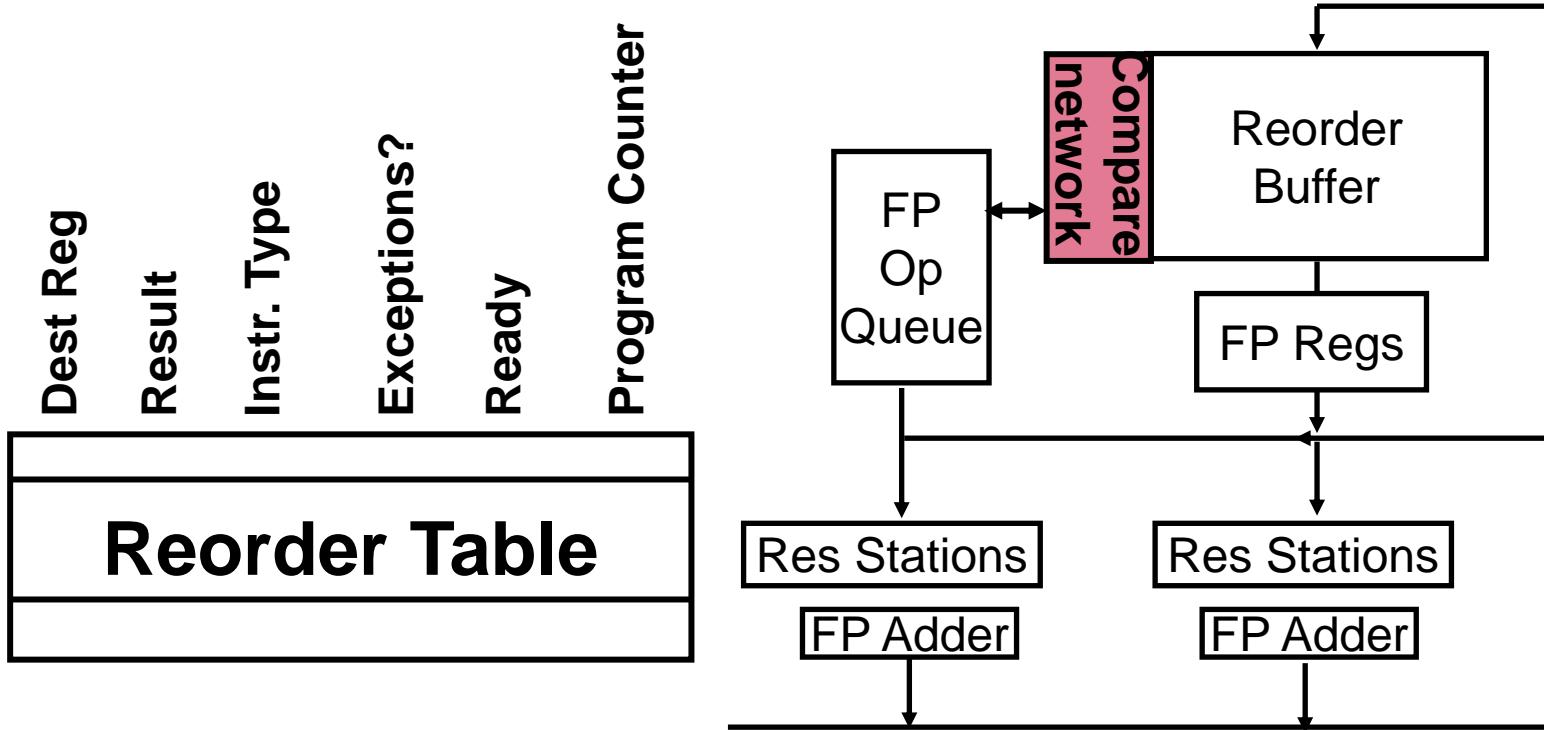
- During the Issue stage of instructions we operate as if as we are predicting that all previous instructions do not generate exceptions
 - Branch prediction, data prediction
 - If we speculate and are wrong, need to back up and restart execution to the point at which we predicted incorrectly
 - This is exactly same as precise exceptions!
- Common technique for precise interrupts/exceptions and speculation: **in-order completion or commit**
 - All modern out-of-order processors typically use a form of re-order buffer (ROB)

HW support for precise interrupts/exceptions

- Idea behind Reorder Buffer (ROB): keep the instructions in a FIFO, with the exact order that they are issued.
 - Each ROB entry contains PC, dest reg/mem, result, exception status
- When an instruction completes execution then the results are placed in the allocated entry in the ROB.
 - Supplies operands to other instruction between execution complete & commit \Rightarrow more registers like RS
 - Tag results with ROB buffer number instead of reservation station
- The instructions change the machine state at **the commit stage** not on the WB \Rightarrow **in order commit** \Rightarrow values at head of ROB are placed in registers
- This technique allows us to cancel/squash speculatively executed instructions during **mispredicted branches or exceptions**

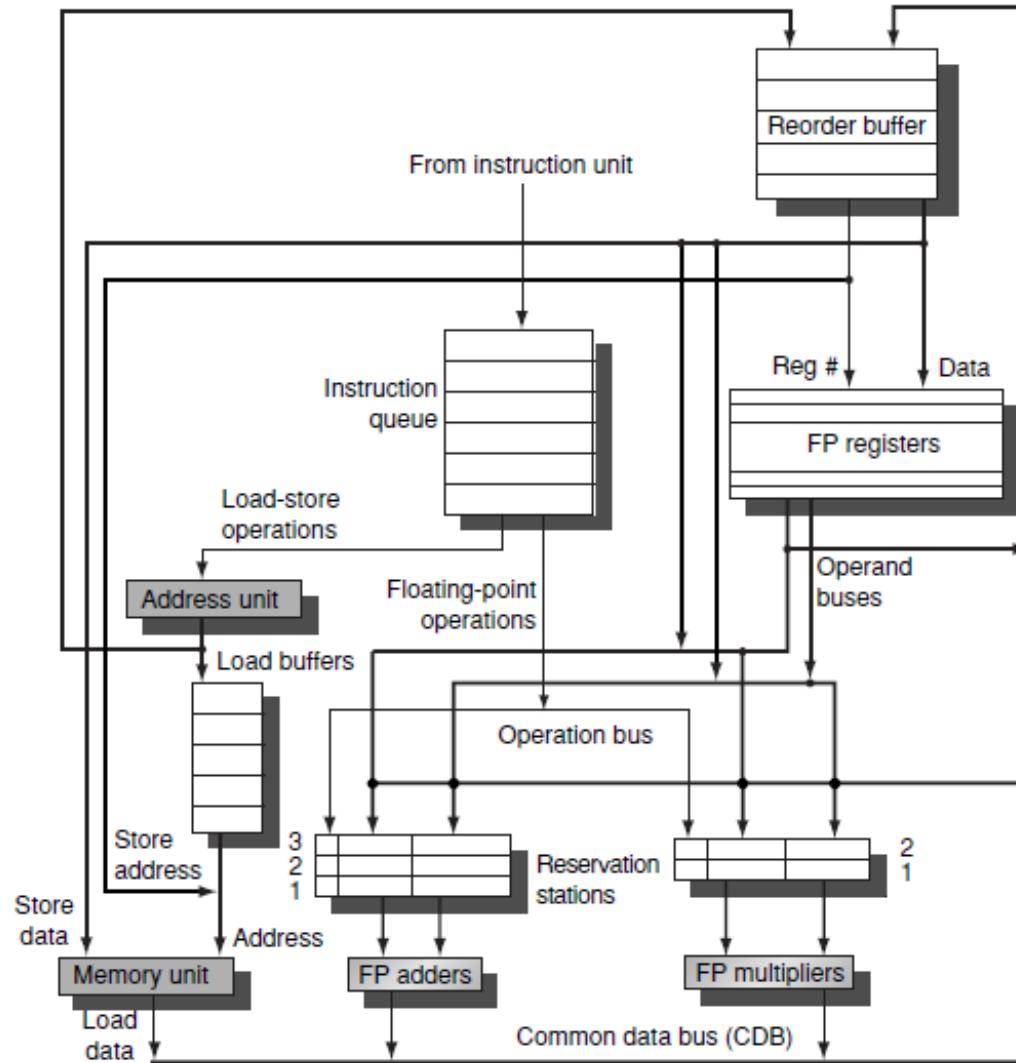


HW Support for Reorder Buffer (ROB)?



- How do we find the last “version” of each register ?
- Multi-ported ROB like the register file
- Integrate store buffer into ROB since we have in order commit. Stores use Result field for ROB tag until data ready on CDB.
- Can we also integrate the reservation stations ?

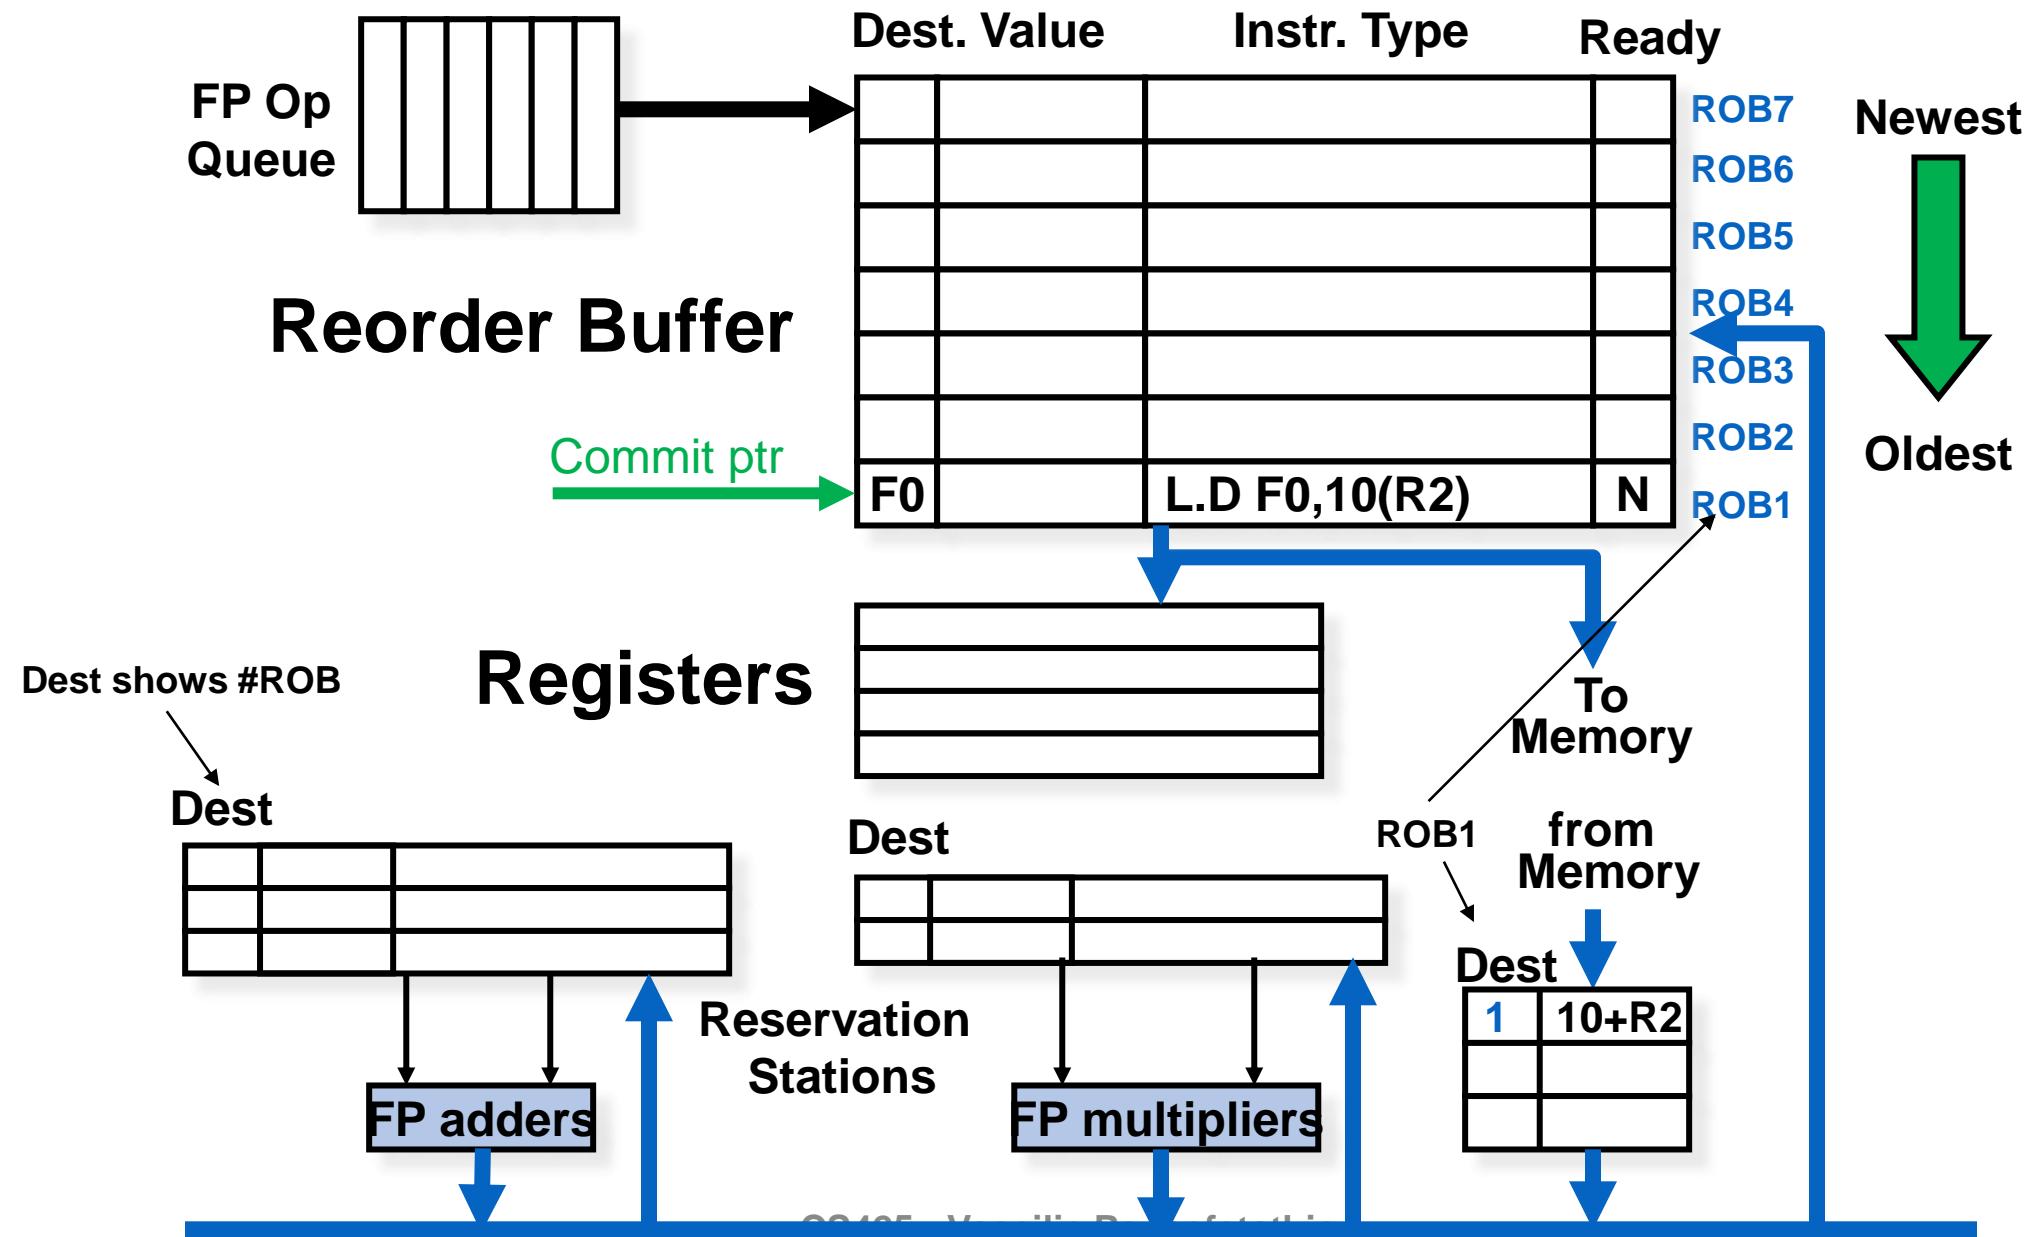
Tomasulo with ROB: Basic Block Diagram



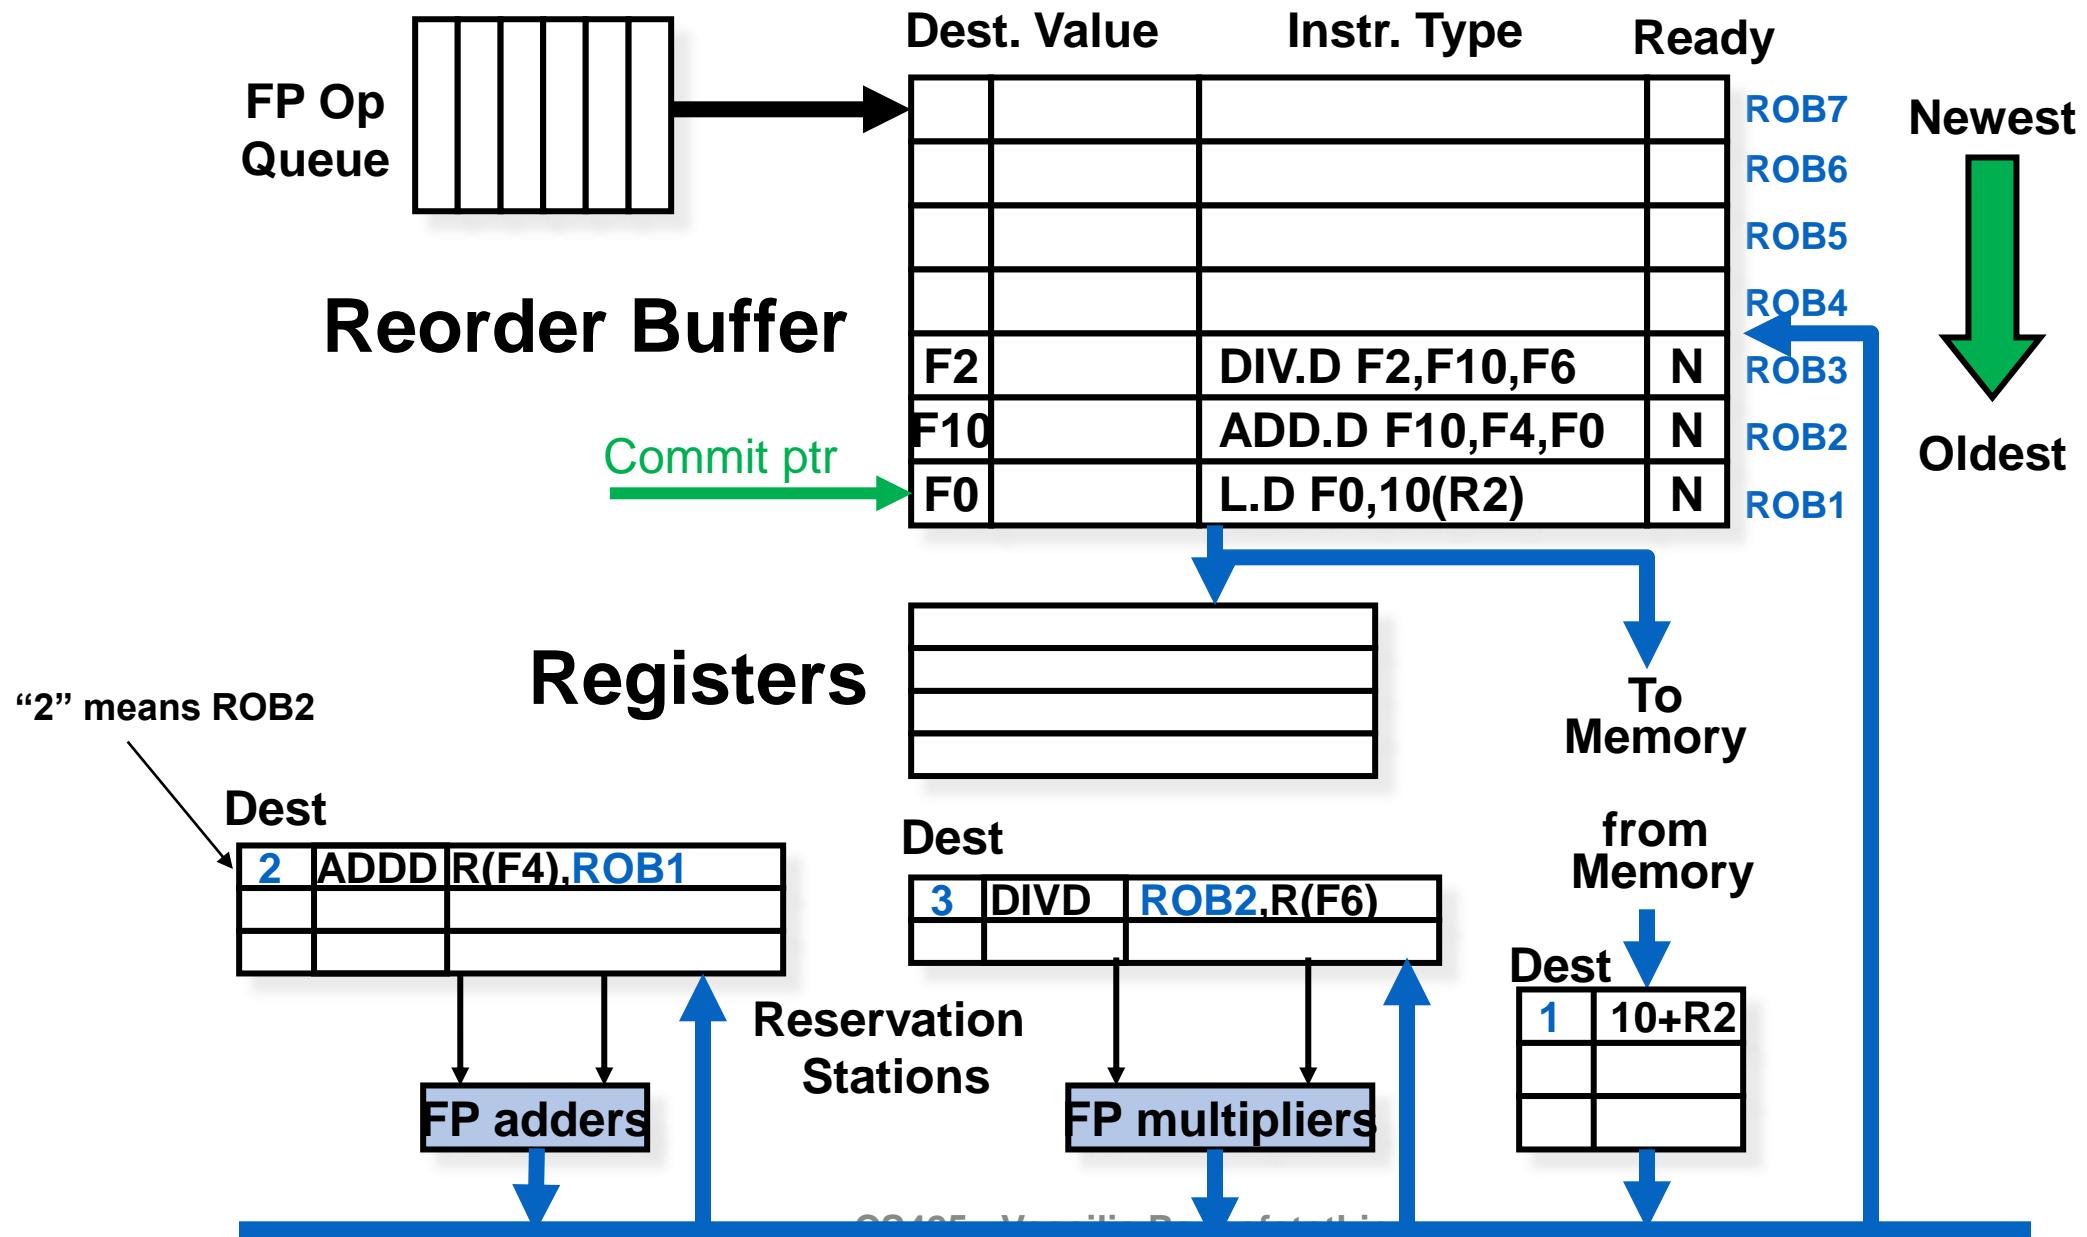
Four Stages of Tomasulo with ROB

1. **Issue:** Get Instruction from Op Queue
 - If there are free reservation stations and **reorder buffer slot**, issue instr & send operands & **reorder buffer no. for destination** (sometimes called “dispatch”)
2. **Execution:** Execute the Instruction in the Execution Unit (EX)
 - When the values of the 2 source regs are ready then execute the instruction; otherwise, watch CDB for result; when both in reservation station, execute; checks RAW (“issue”)
3. **Write result:** End of Execution (WB)
 - Write on Common Data Bus to all awaiting FUs & **reorder buffer**; mark reservation station available.
4. **Commit:** Update the dst reg with the value from the **reorder buffer**
 - When instr. at head of reorder buffer & result present, update register with result (or store to memory) and remove instr from reorder buffer. Mispredicted branch or exception flushes reorder buffer. (also called “graduation” or “retirement”)

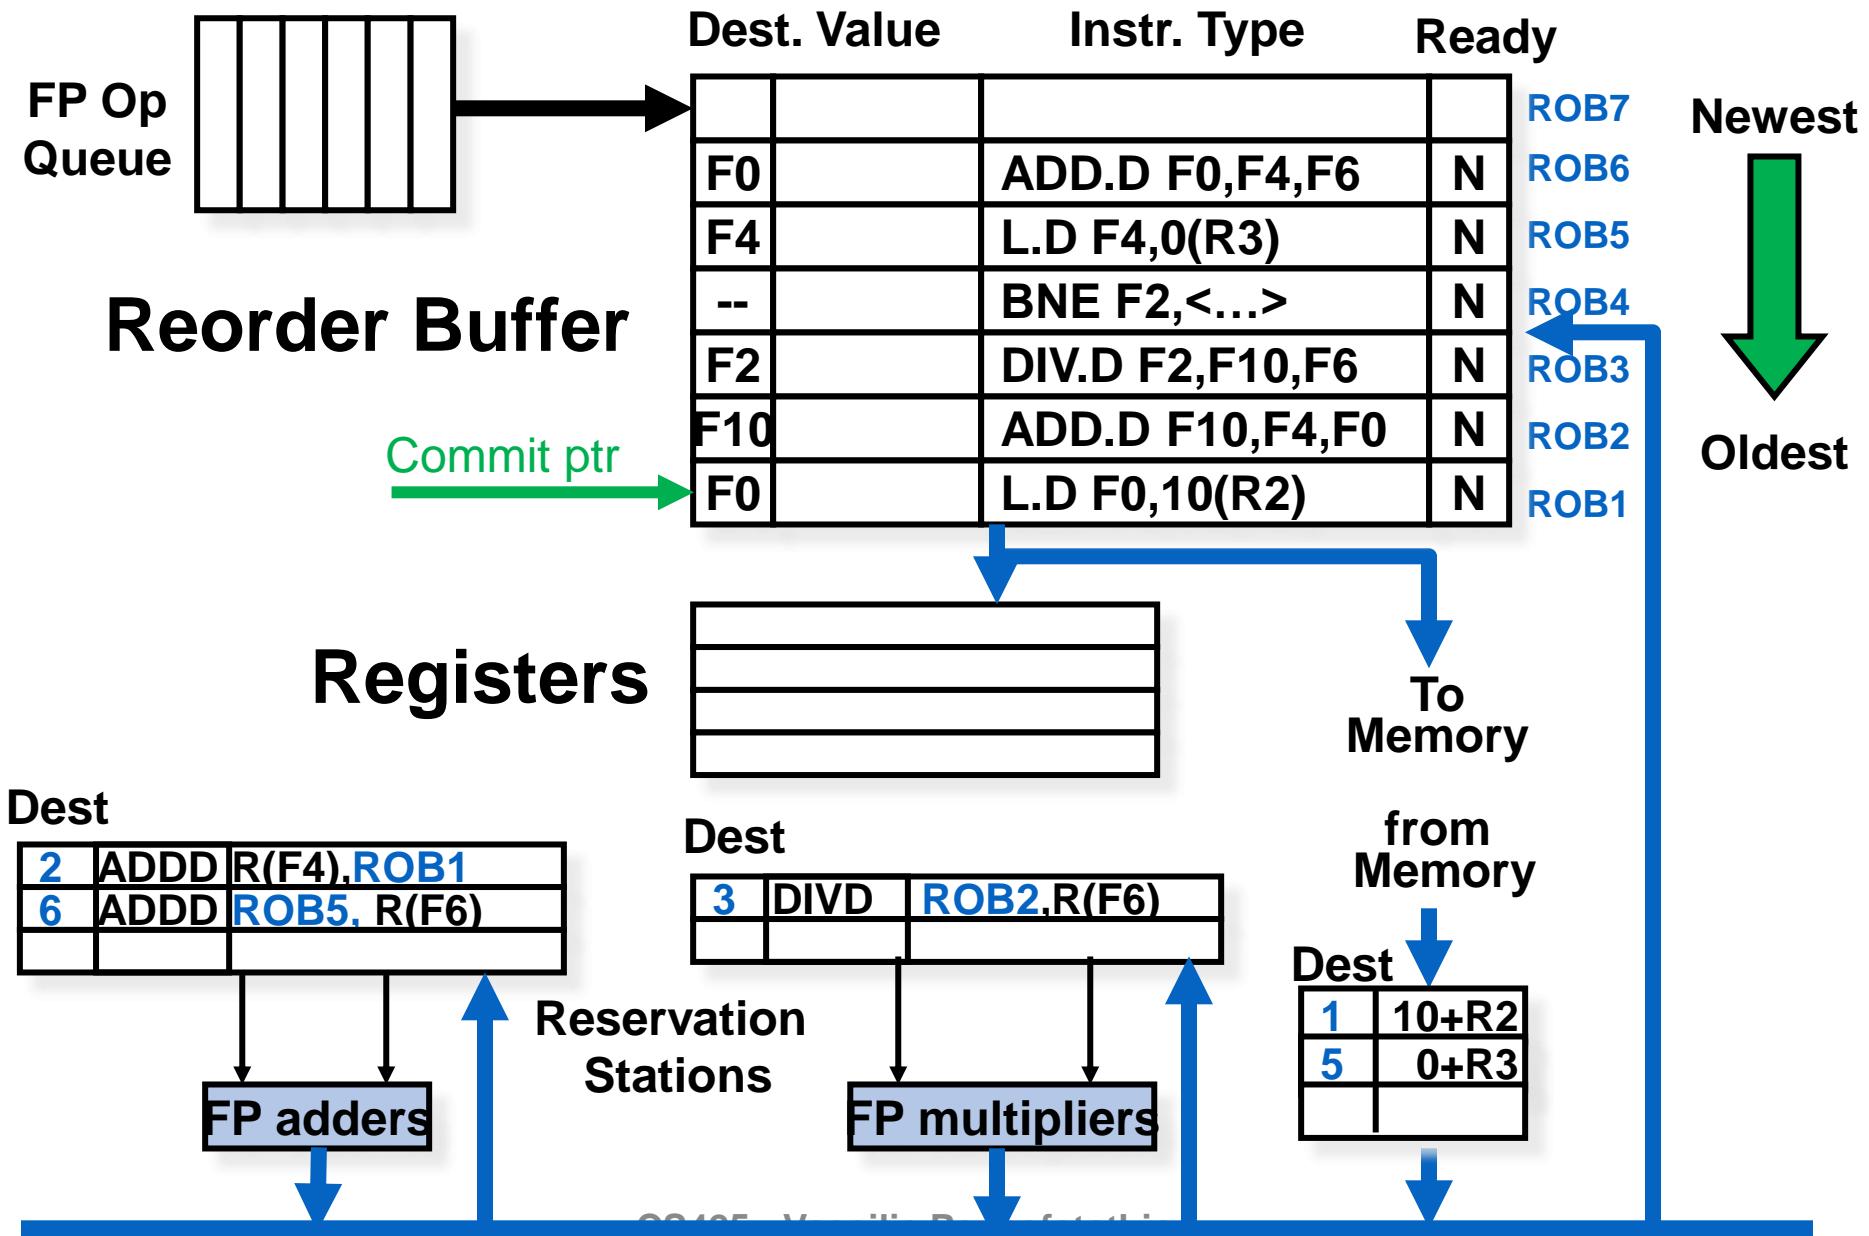
Tomasulo With Reorder buffer



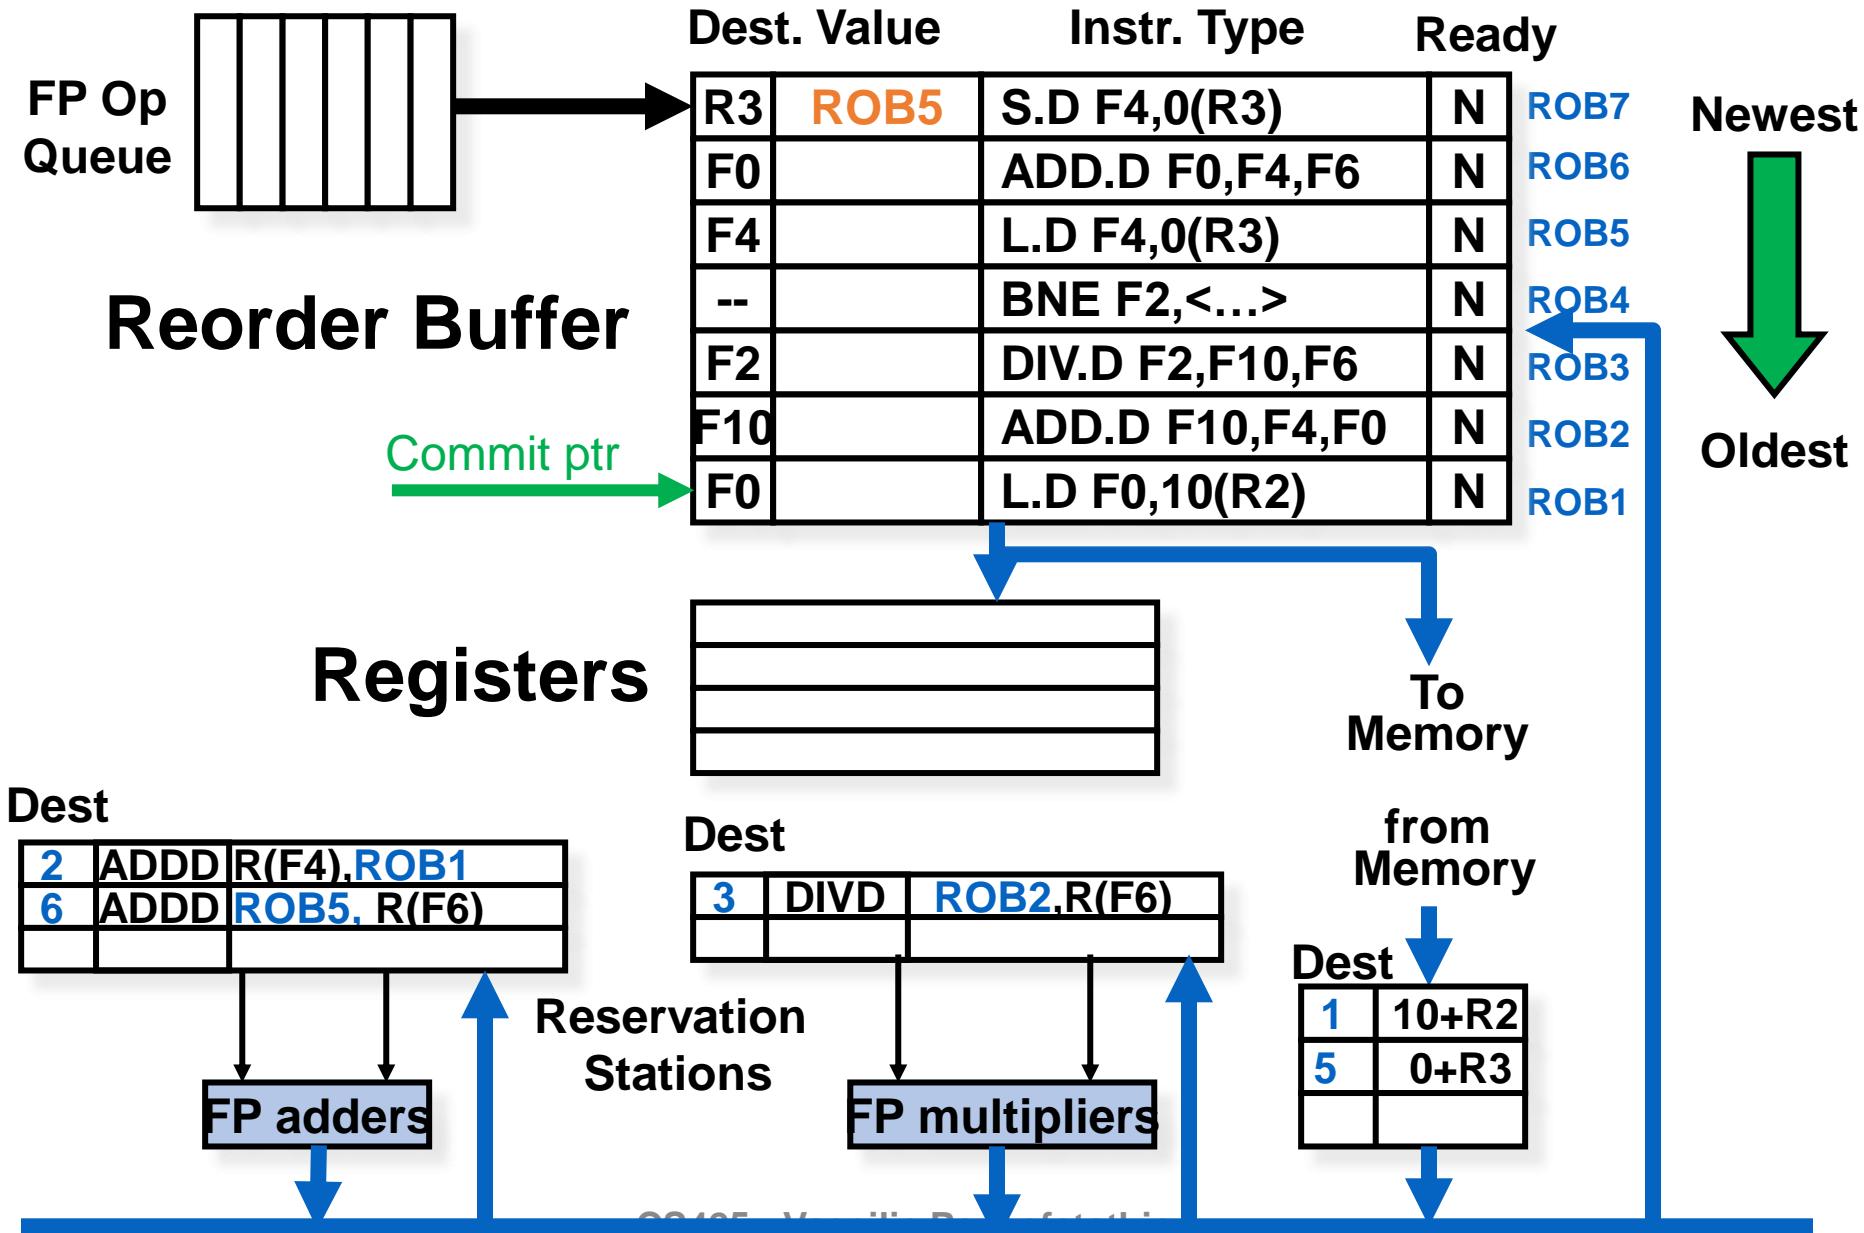
Reorder buffer (after 2 cycles)



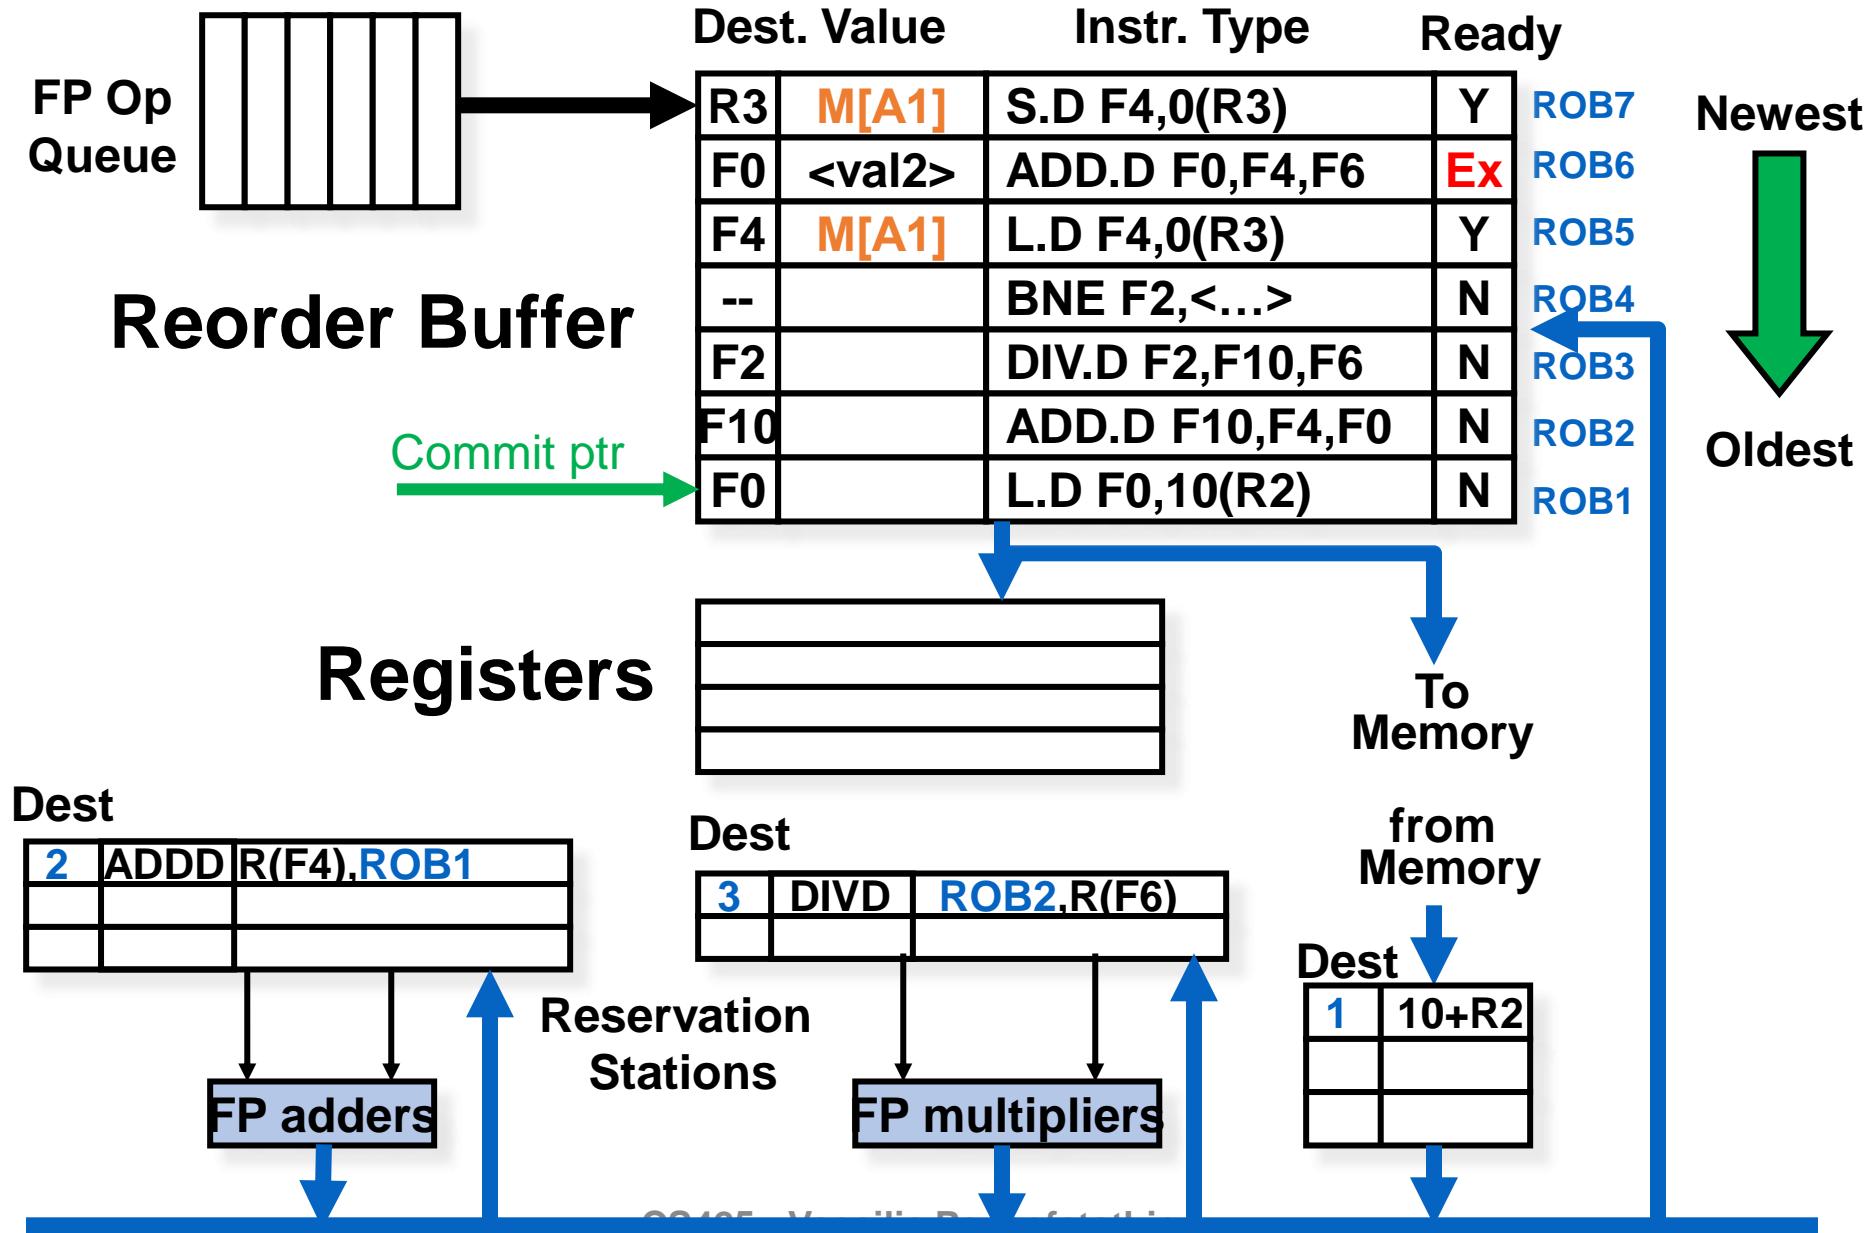
Reorder buffer (after 3 cycles)



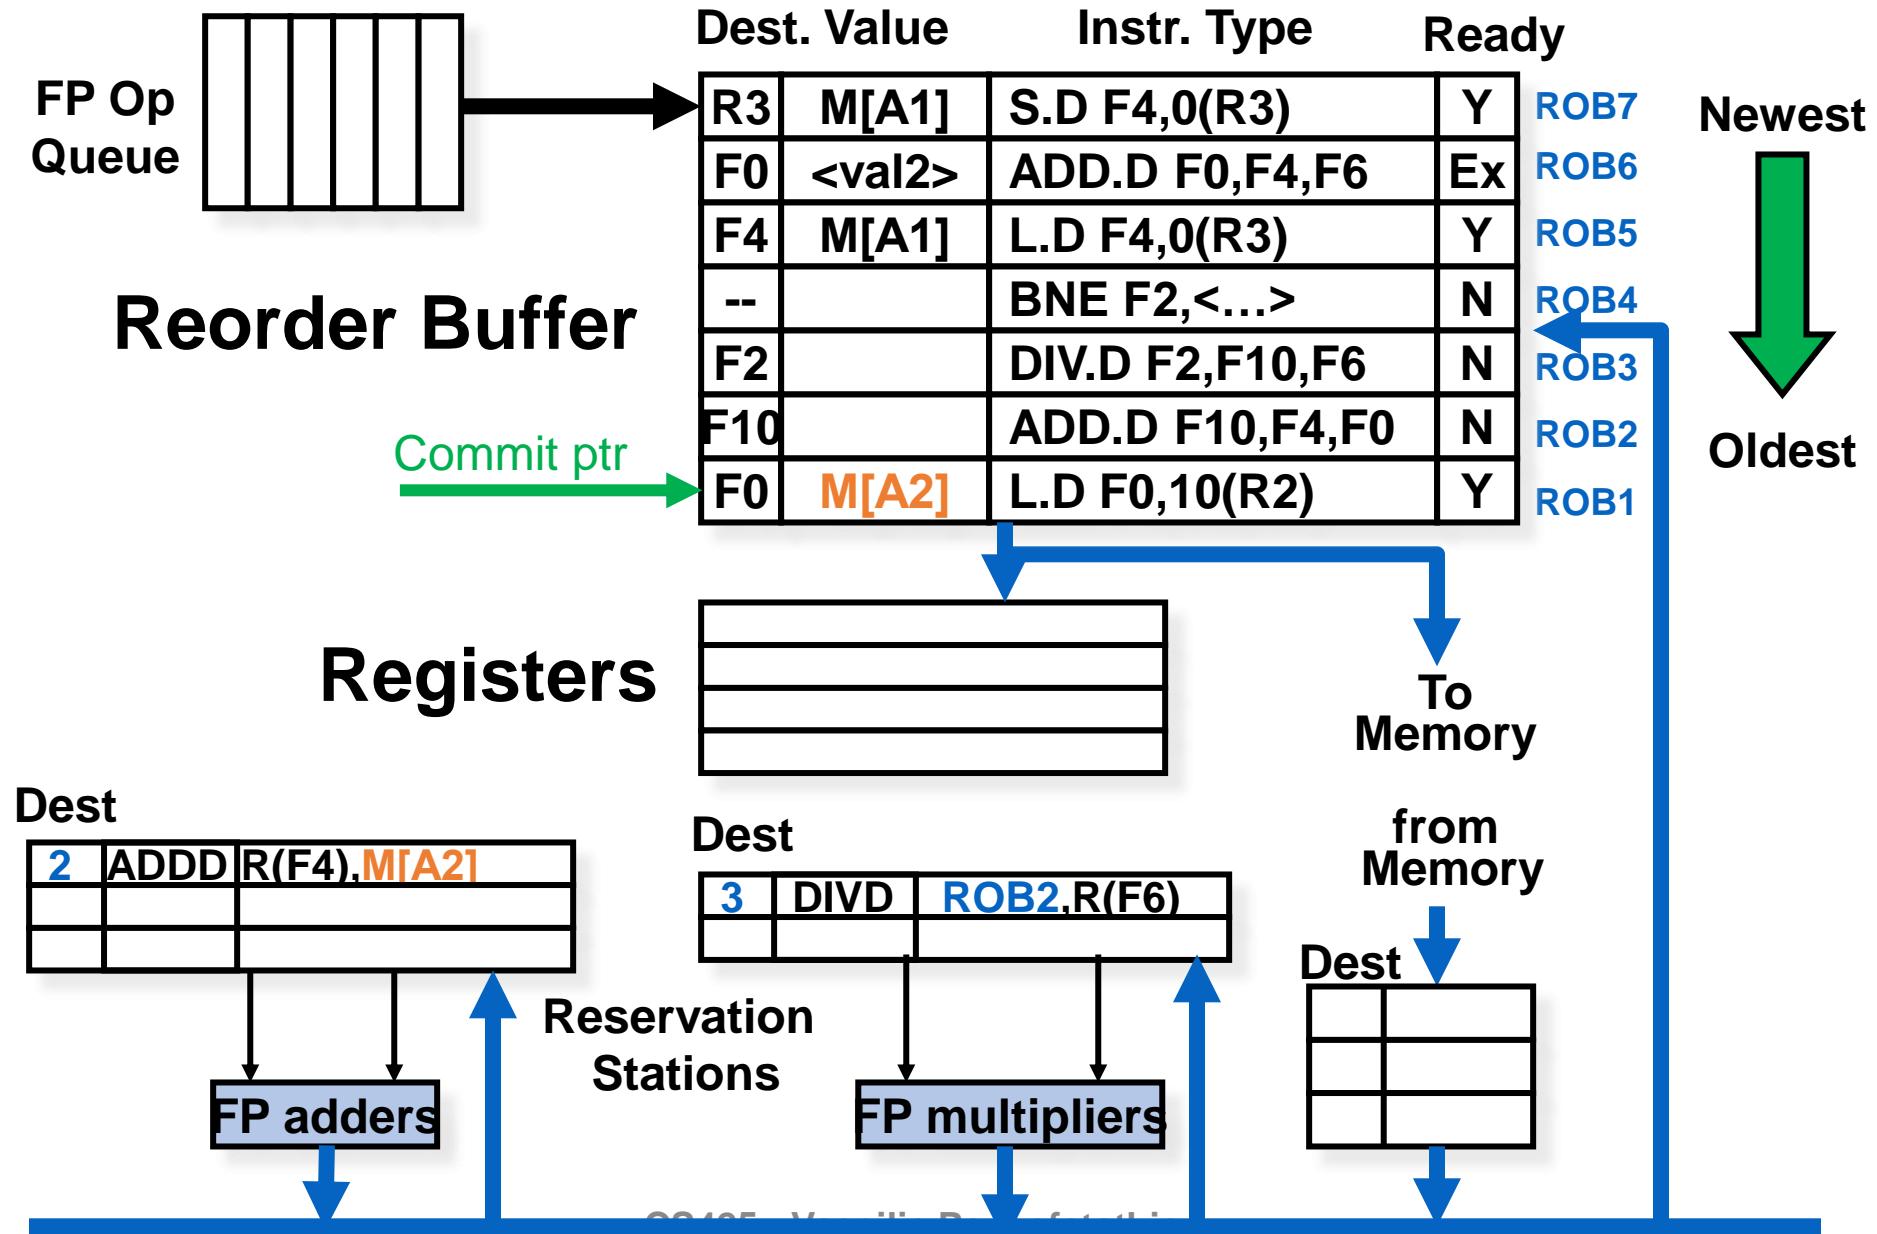
Reorder buffer (after 1 cycle)



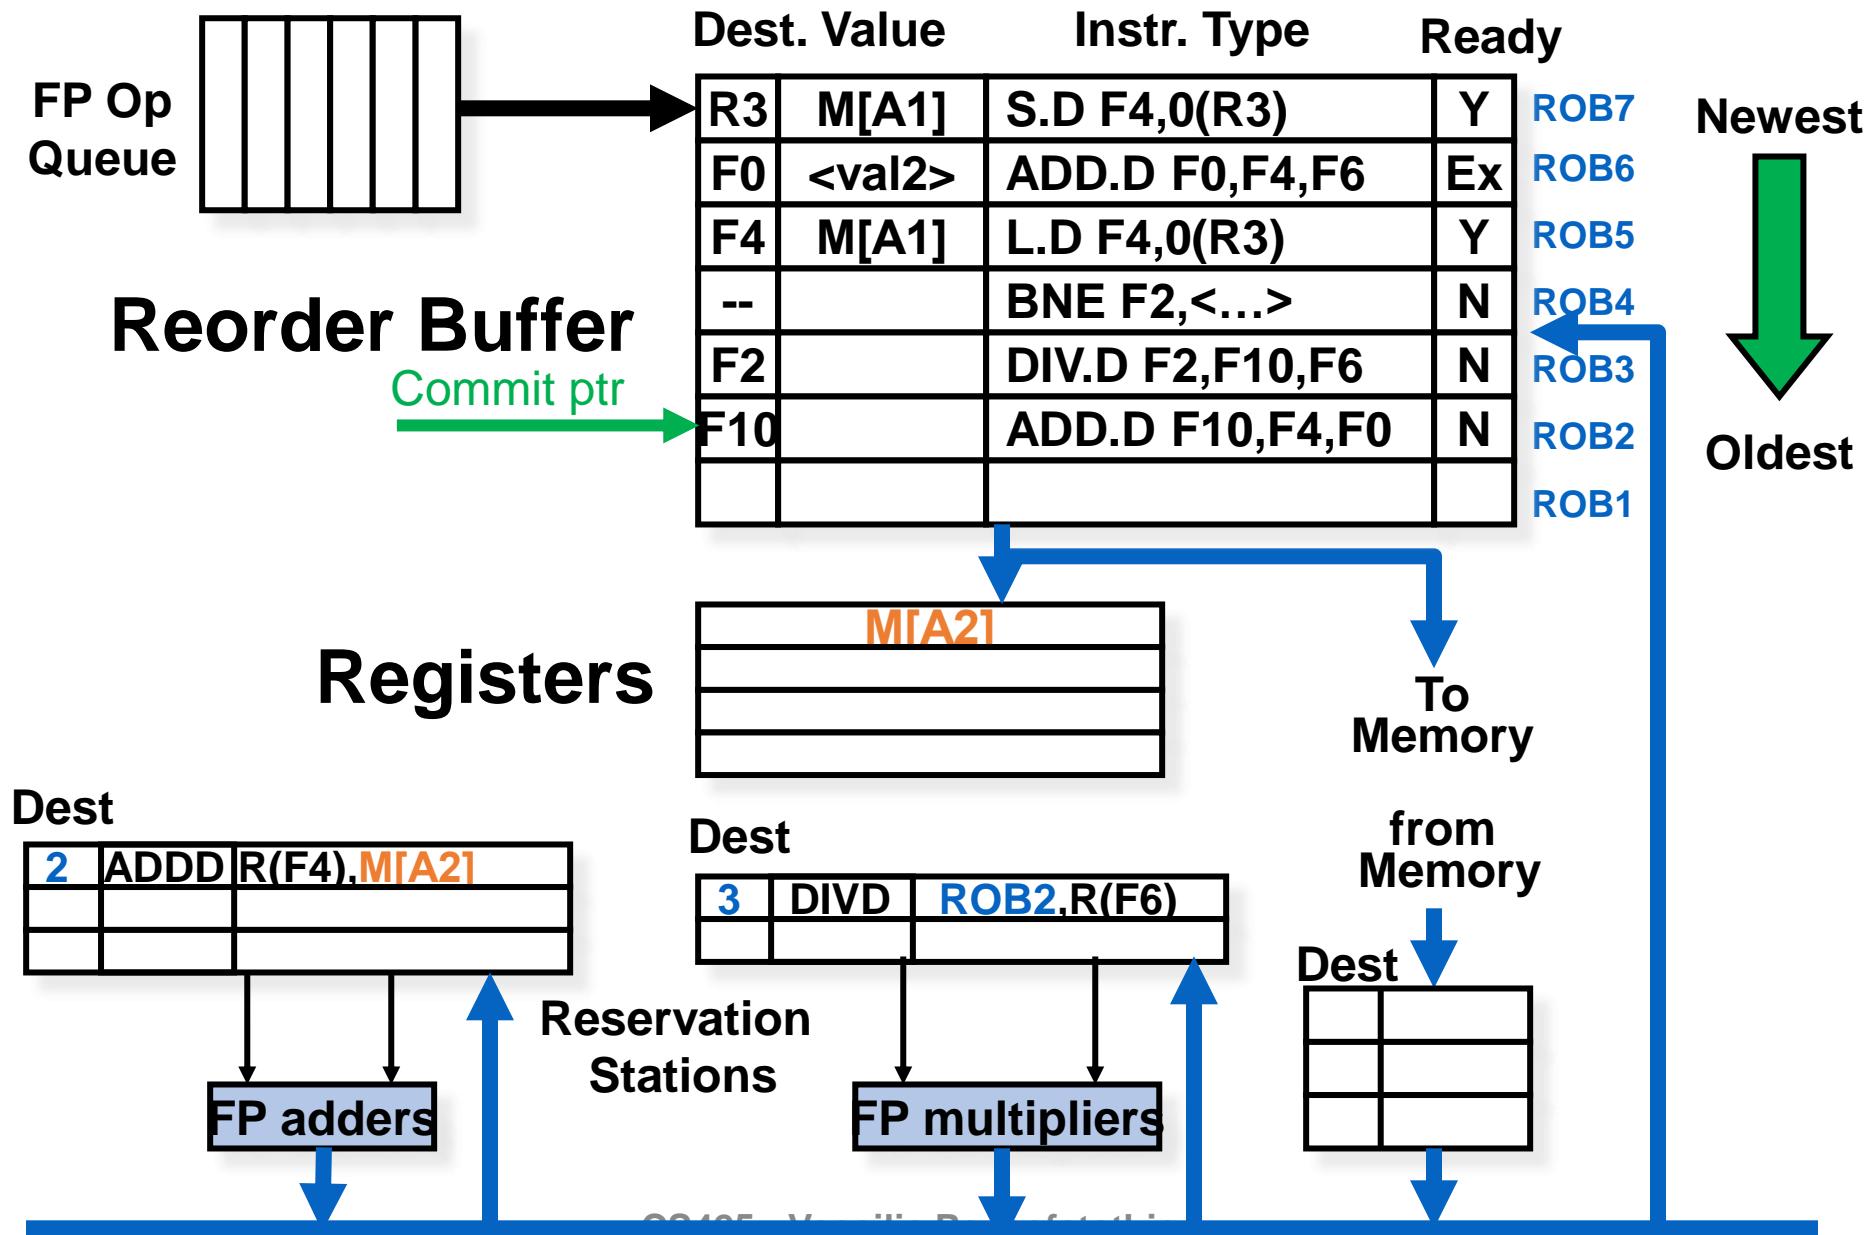
Tomasulo With Reorder buffer



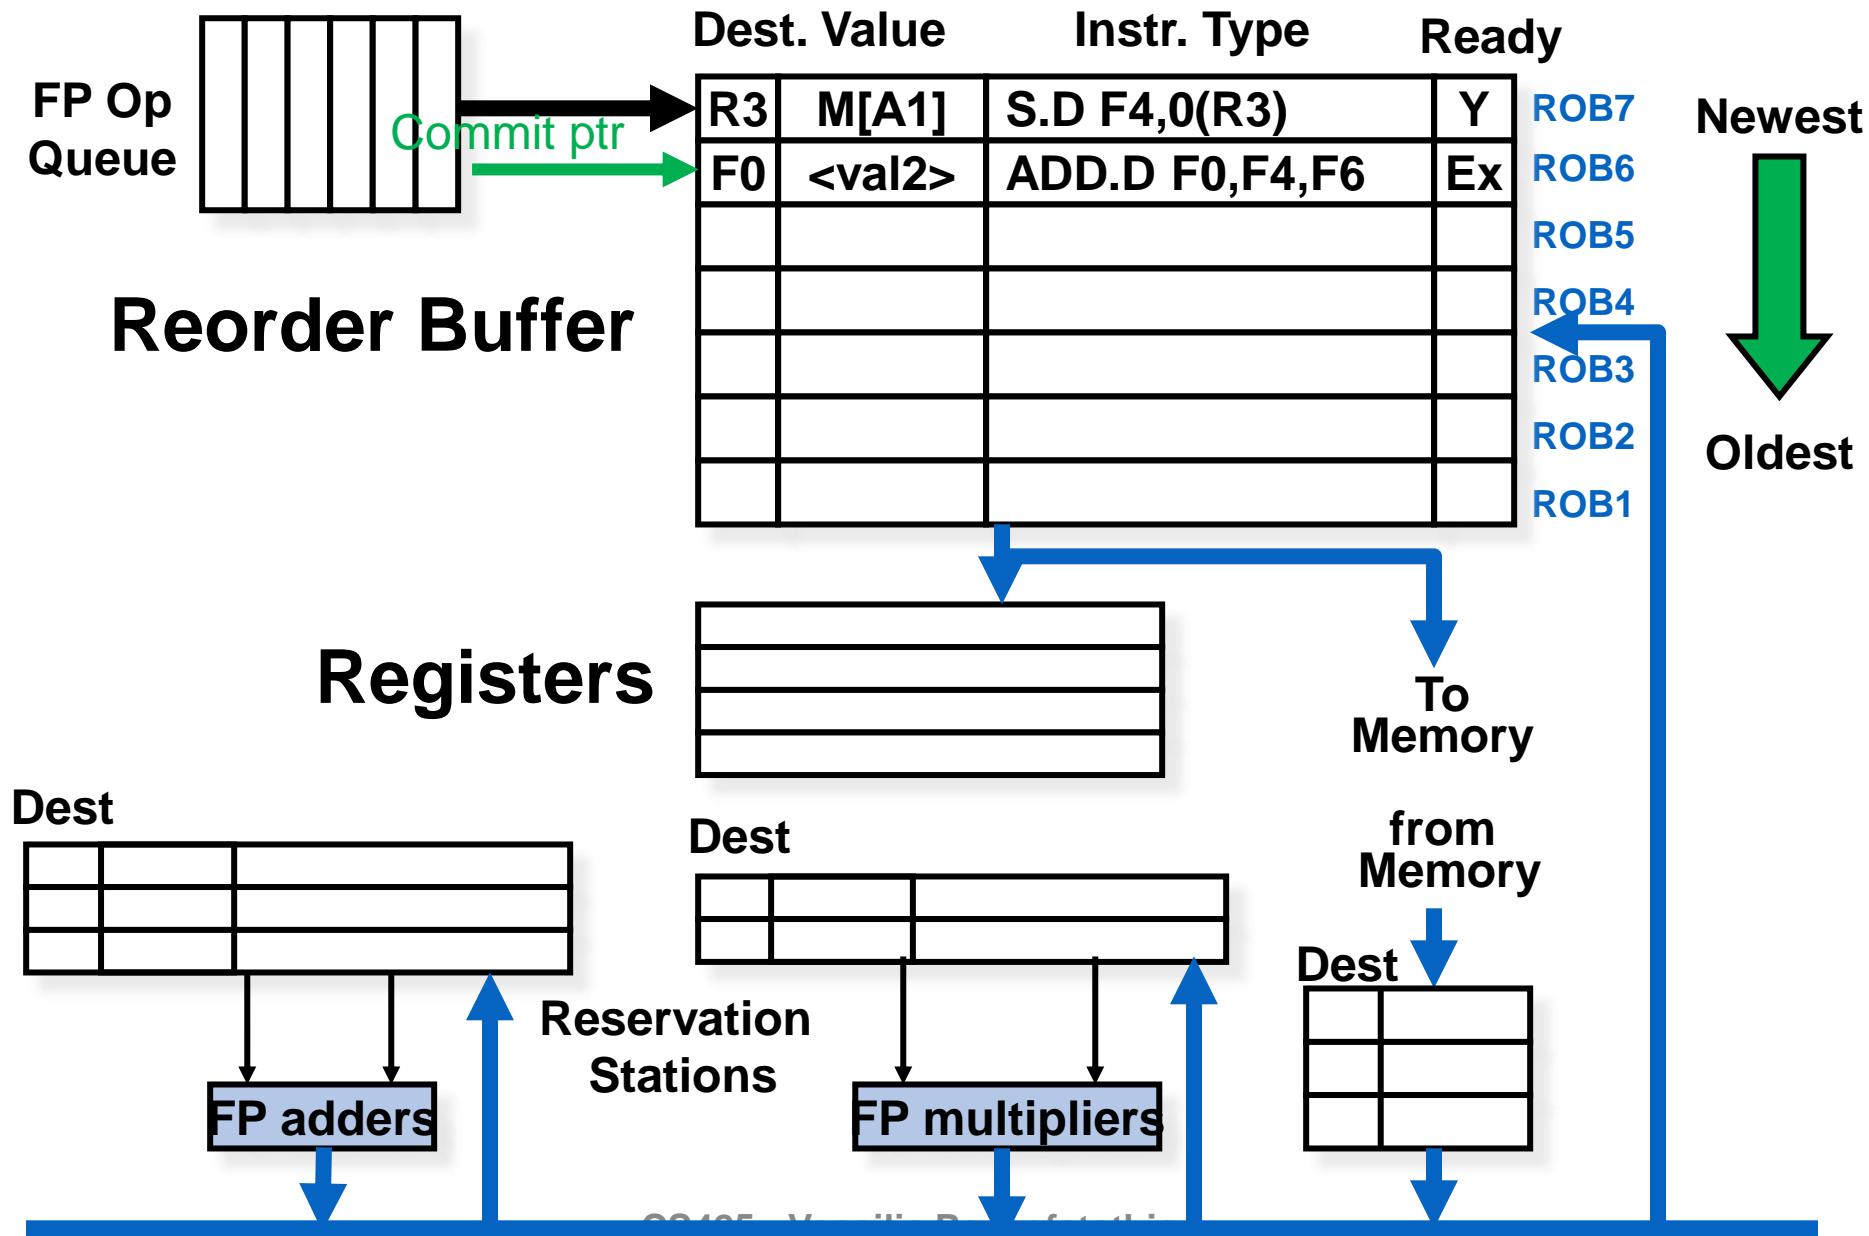
Tomasulo With Reorder buffer



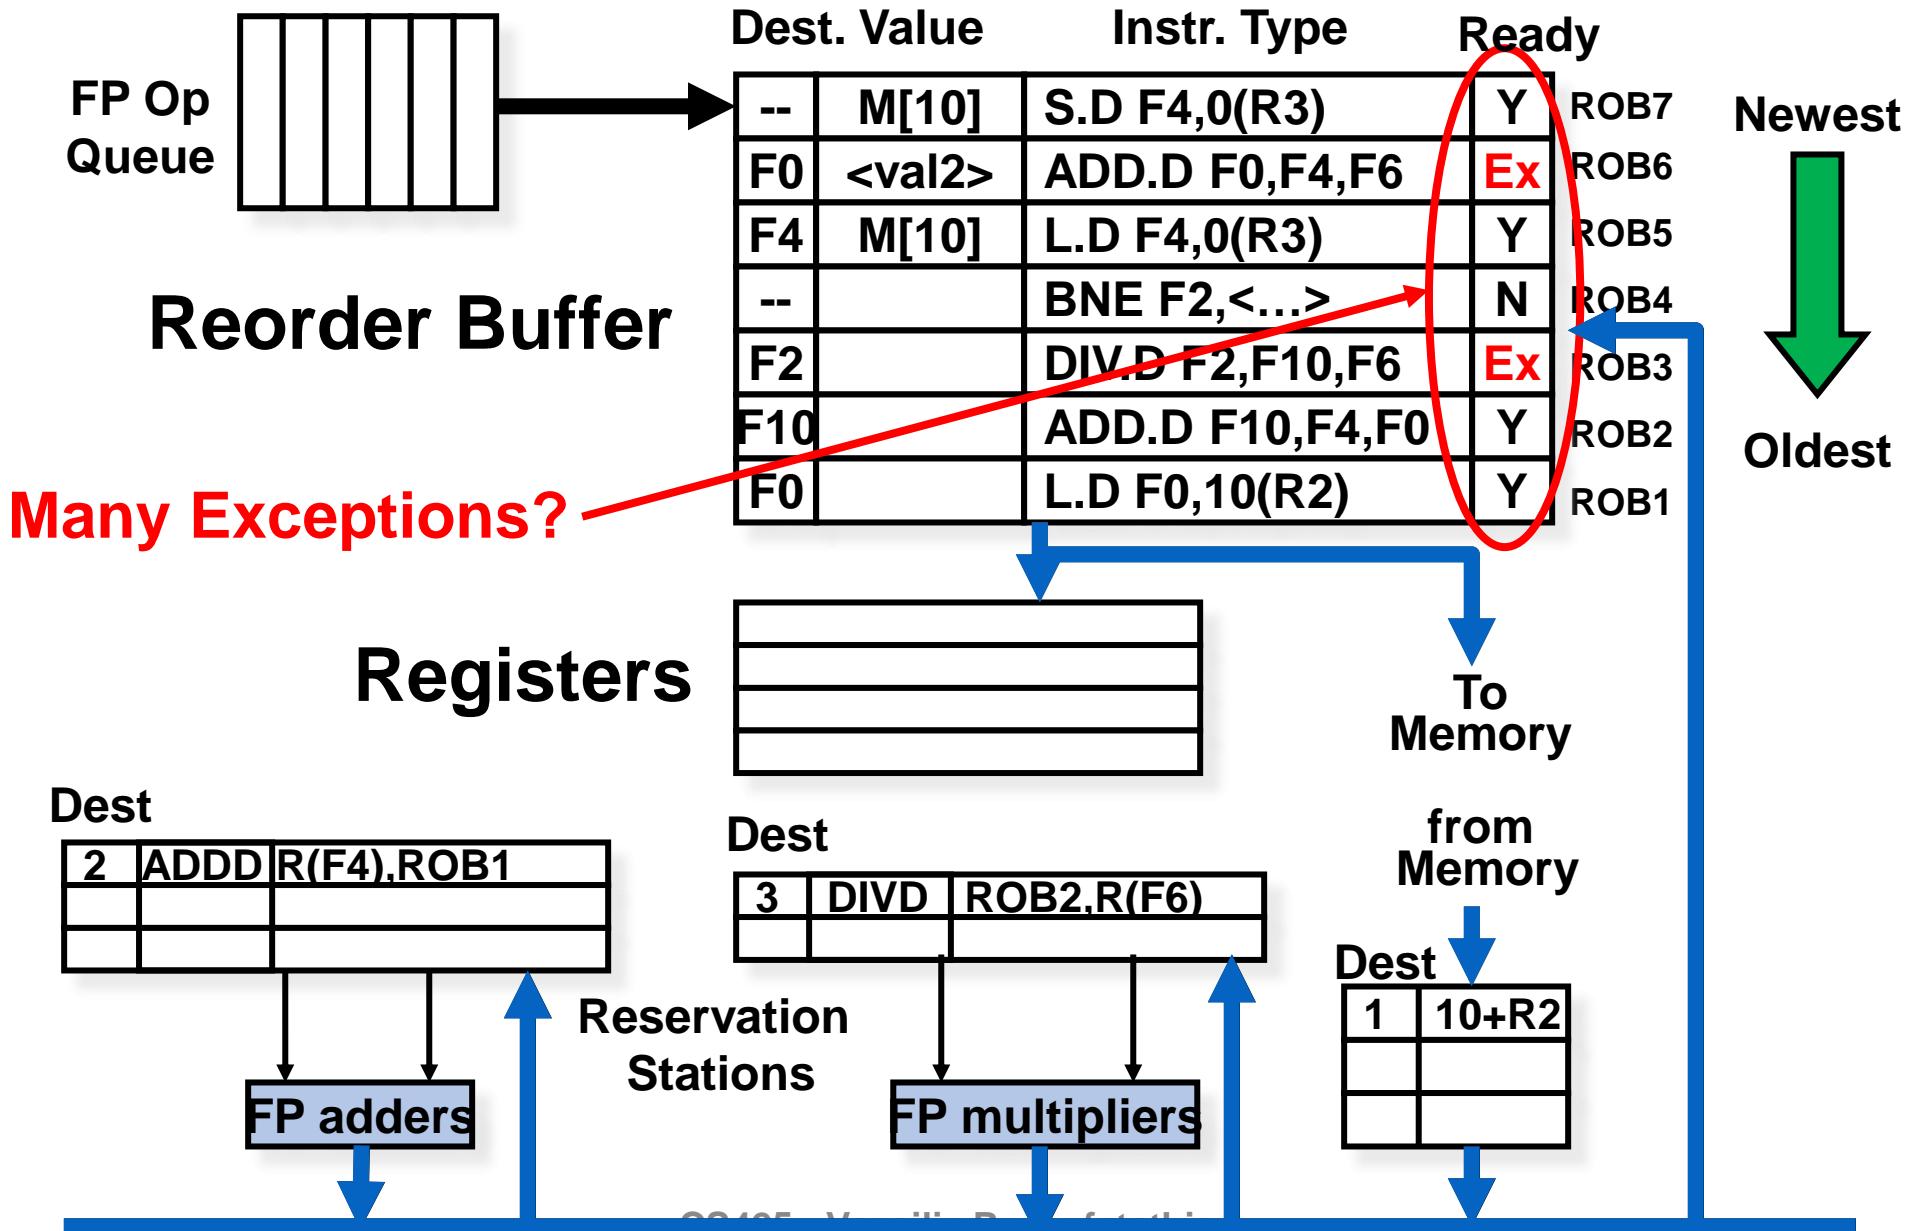
Tomasulo With Reorder buffer



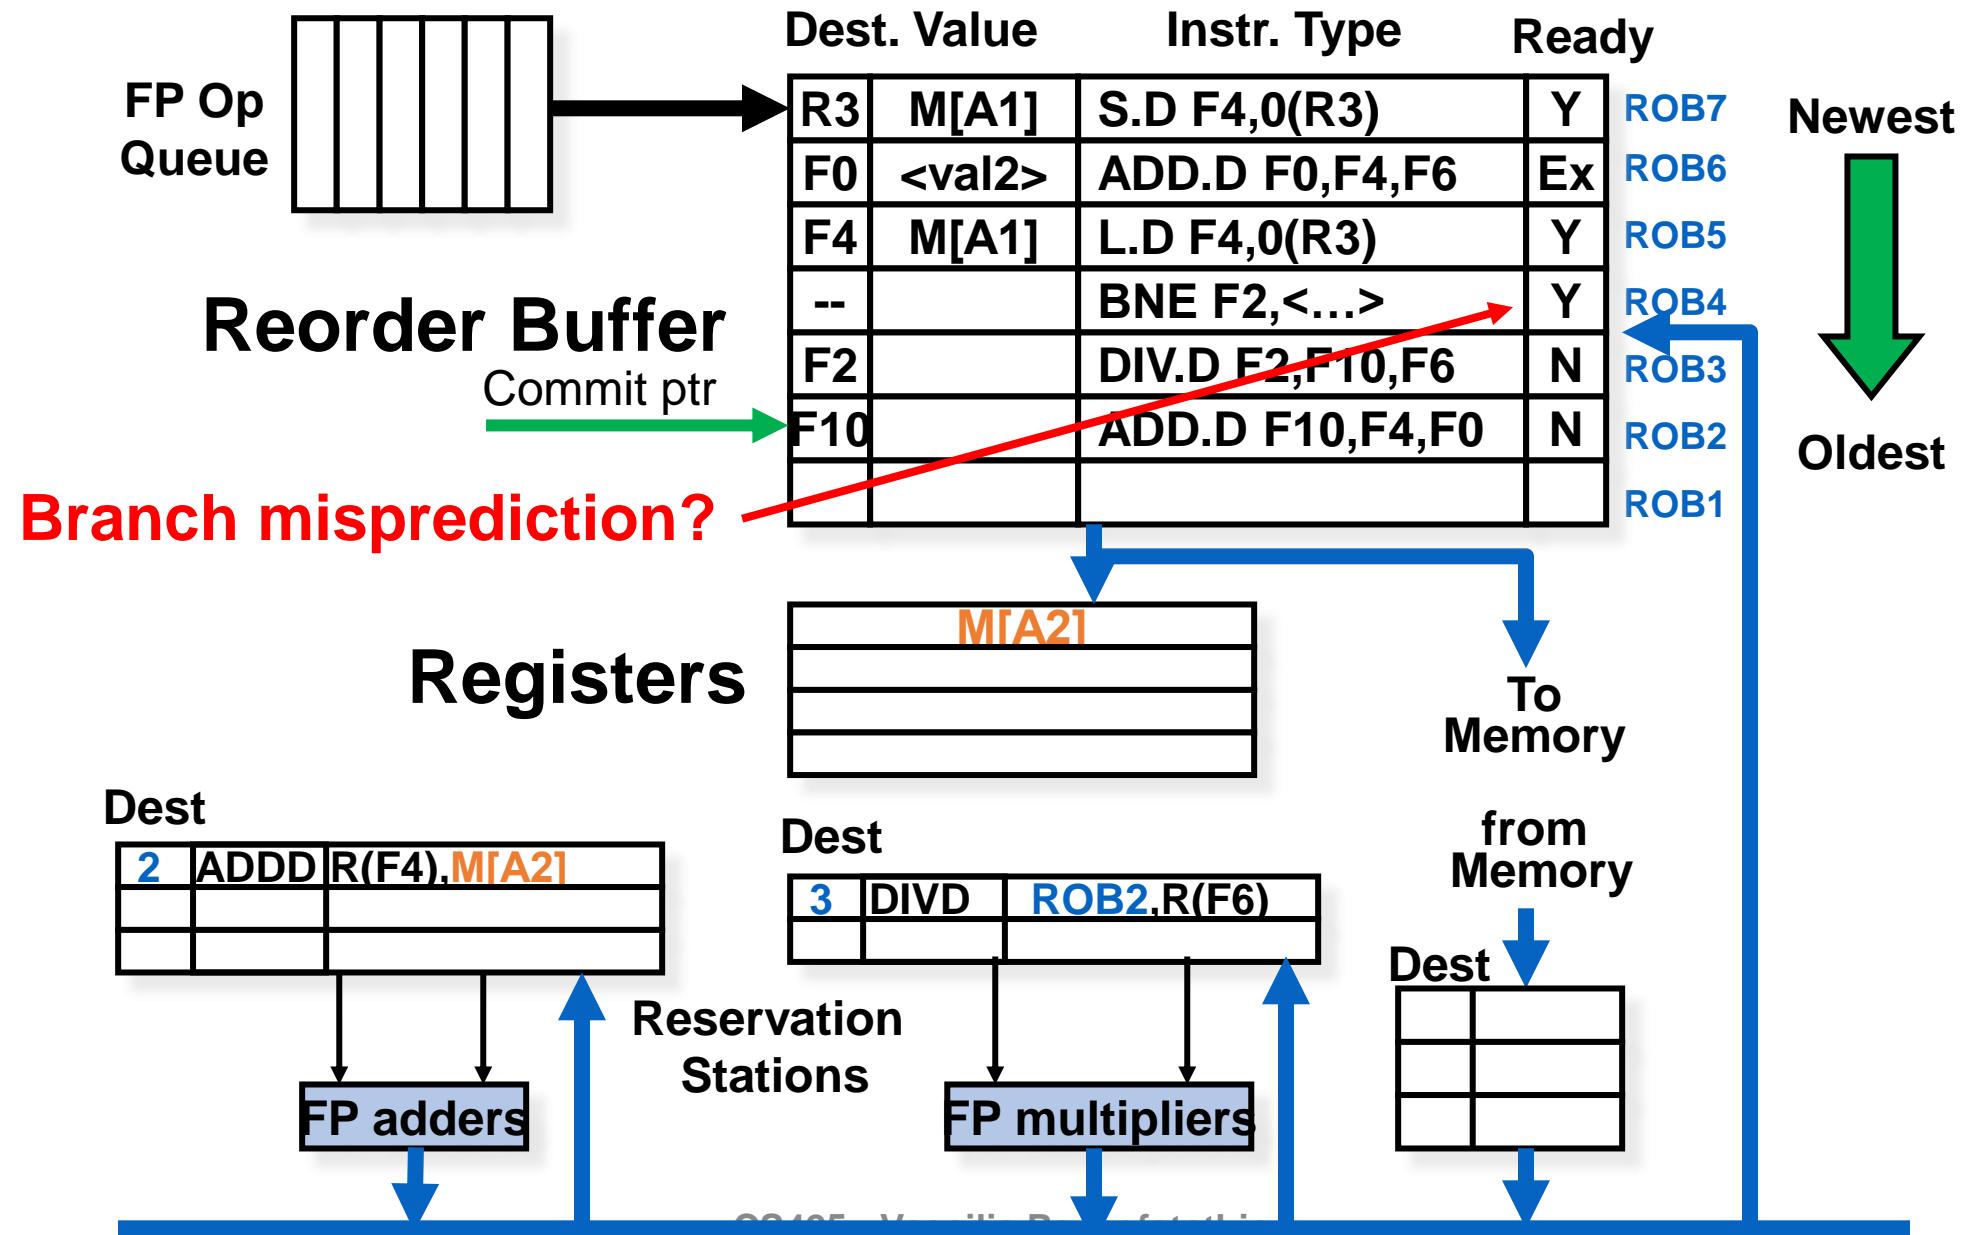
Tomasulo With Reorder buffer



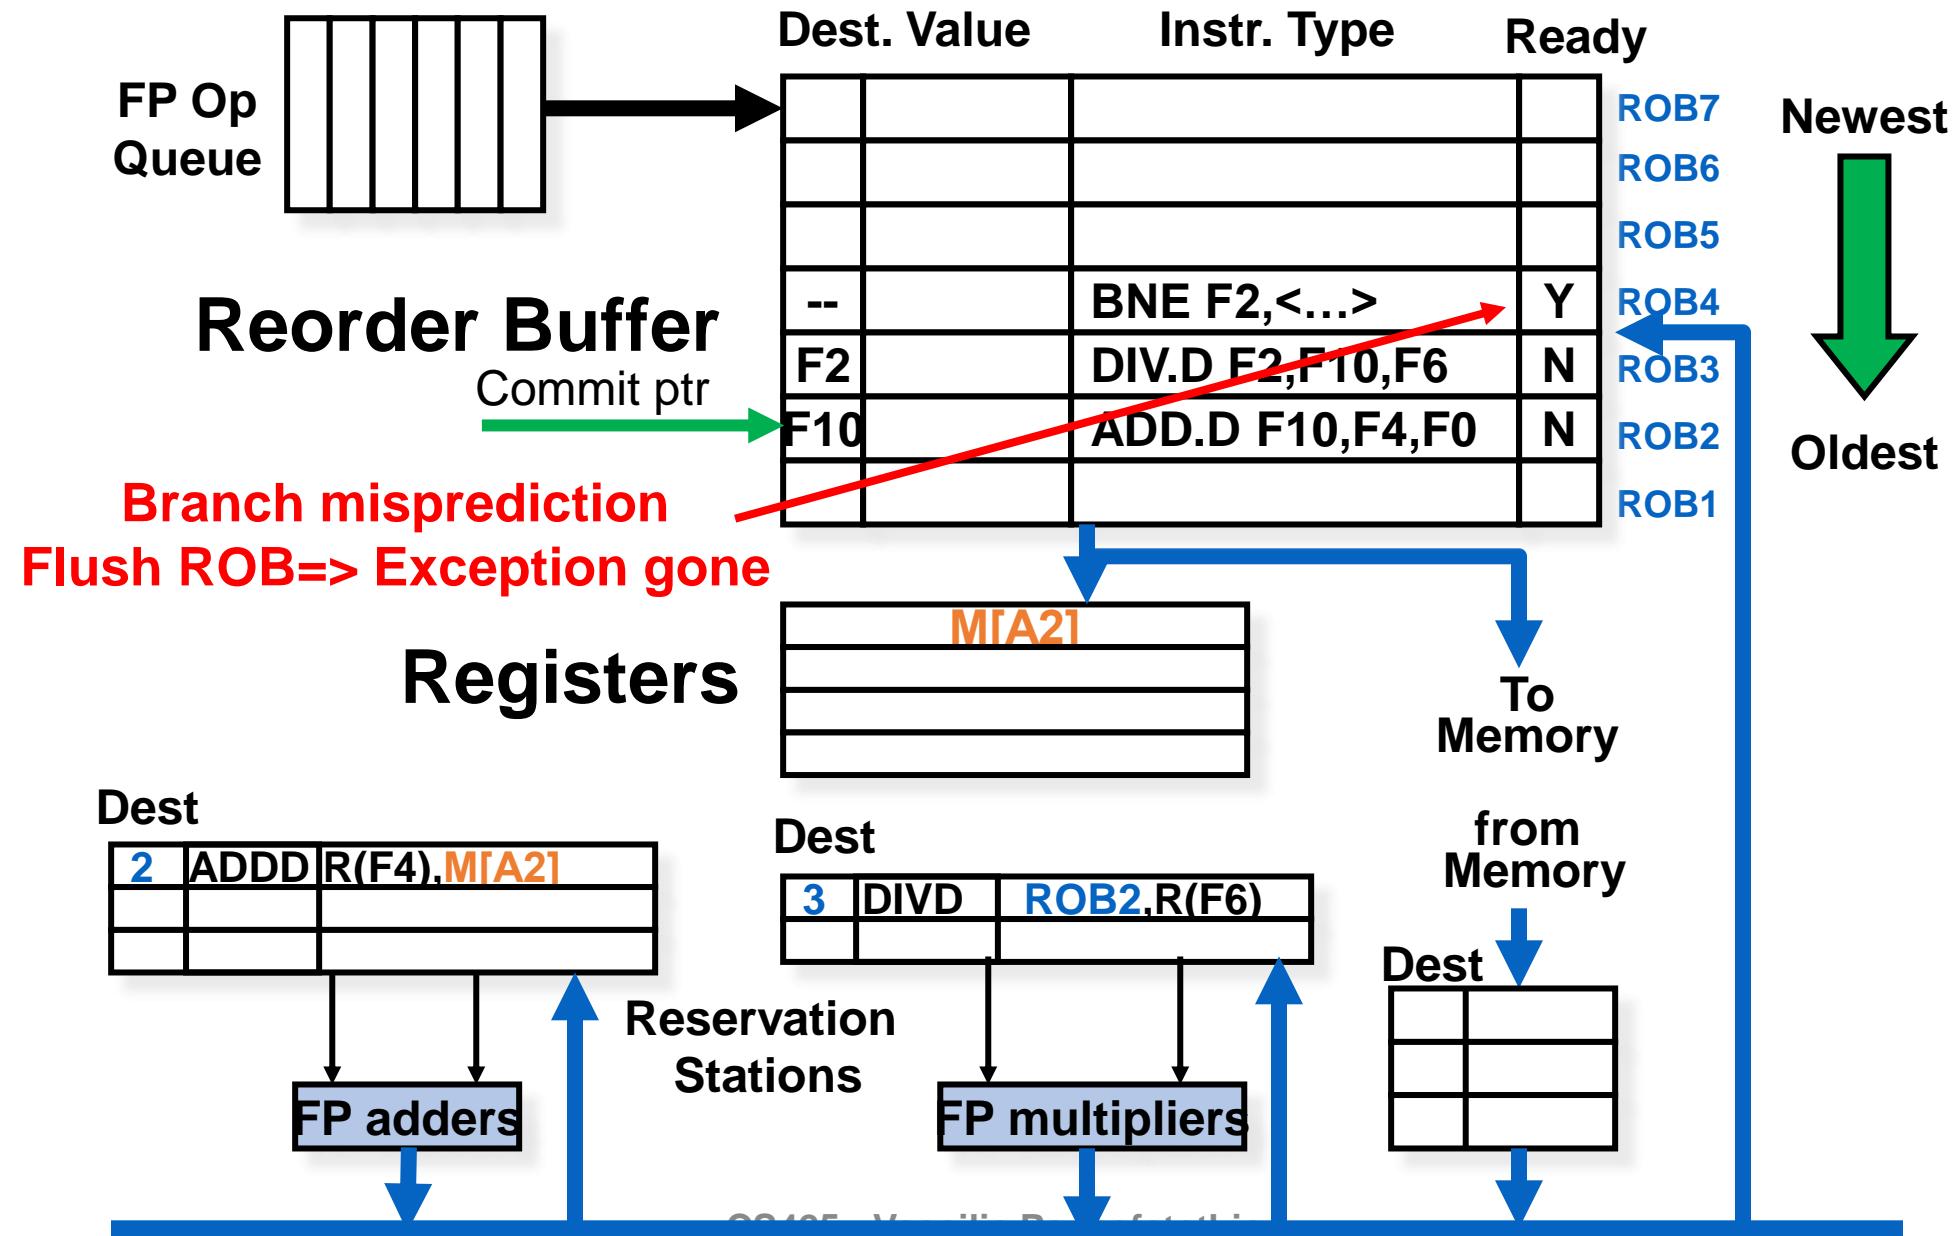
Reorder buffer: Precise Exceptions



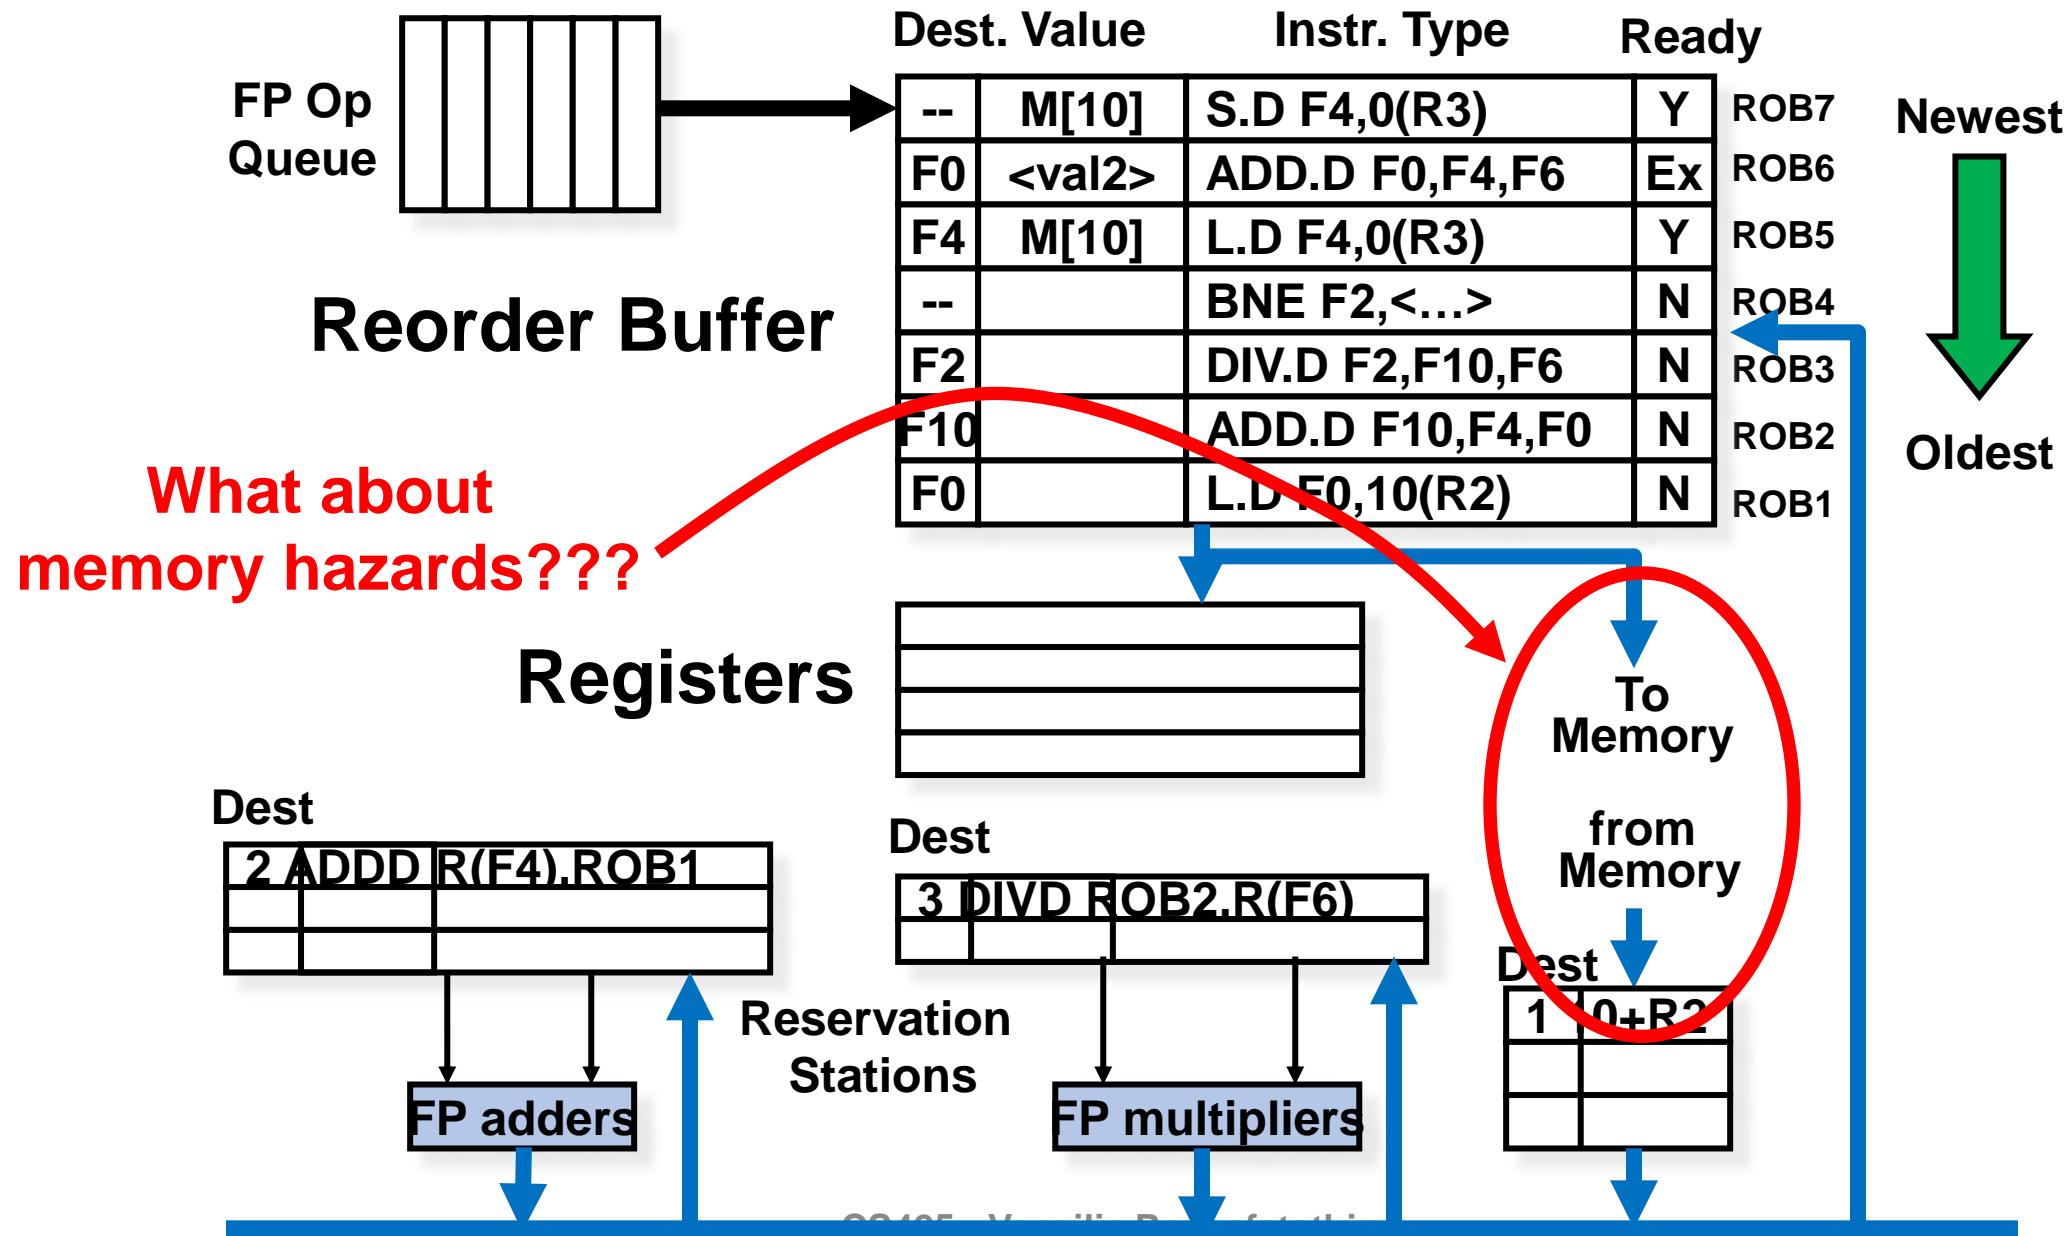
Reorder buffer: Branch Misprediction



Reorder buffer: Branch Misprediction



Tomasulo With Reorder buffer



Memory Disambiguation: WAW/WAR Hazards

- Like Hazards in Register File, we must avoid hazards through memory:
 - WAW and WAR hazards through memory are eliminated with speculation because **the actual updating of memory occurs in order**, when a store is at the head of the ROB, and hence, no earlier loads or stores can still be pending.

Memory Disambiguation: RAW Hazards

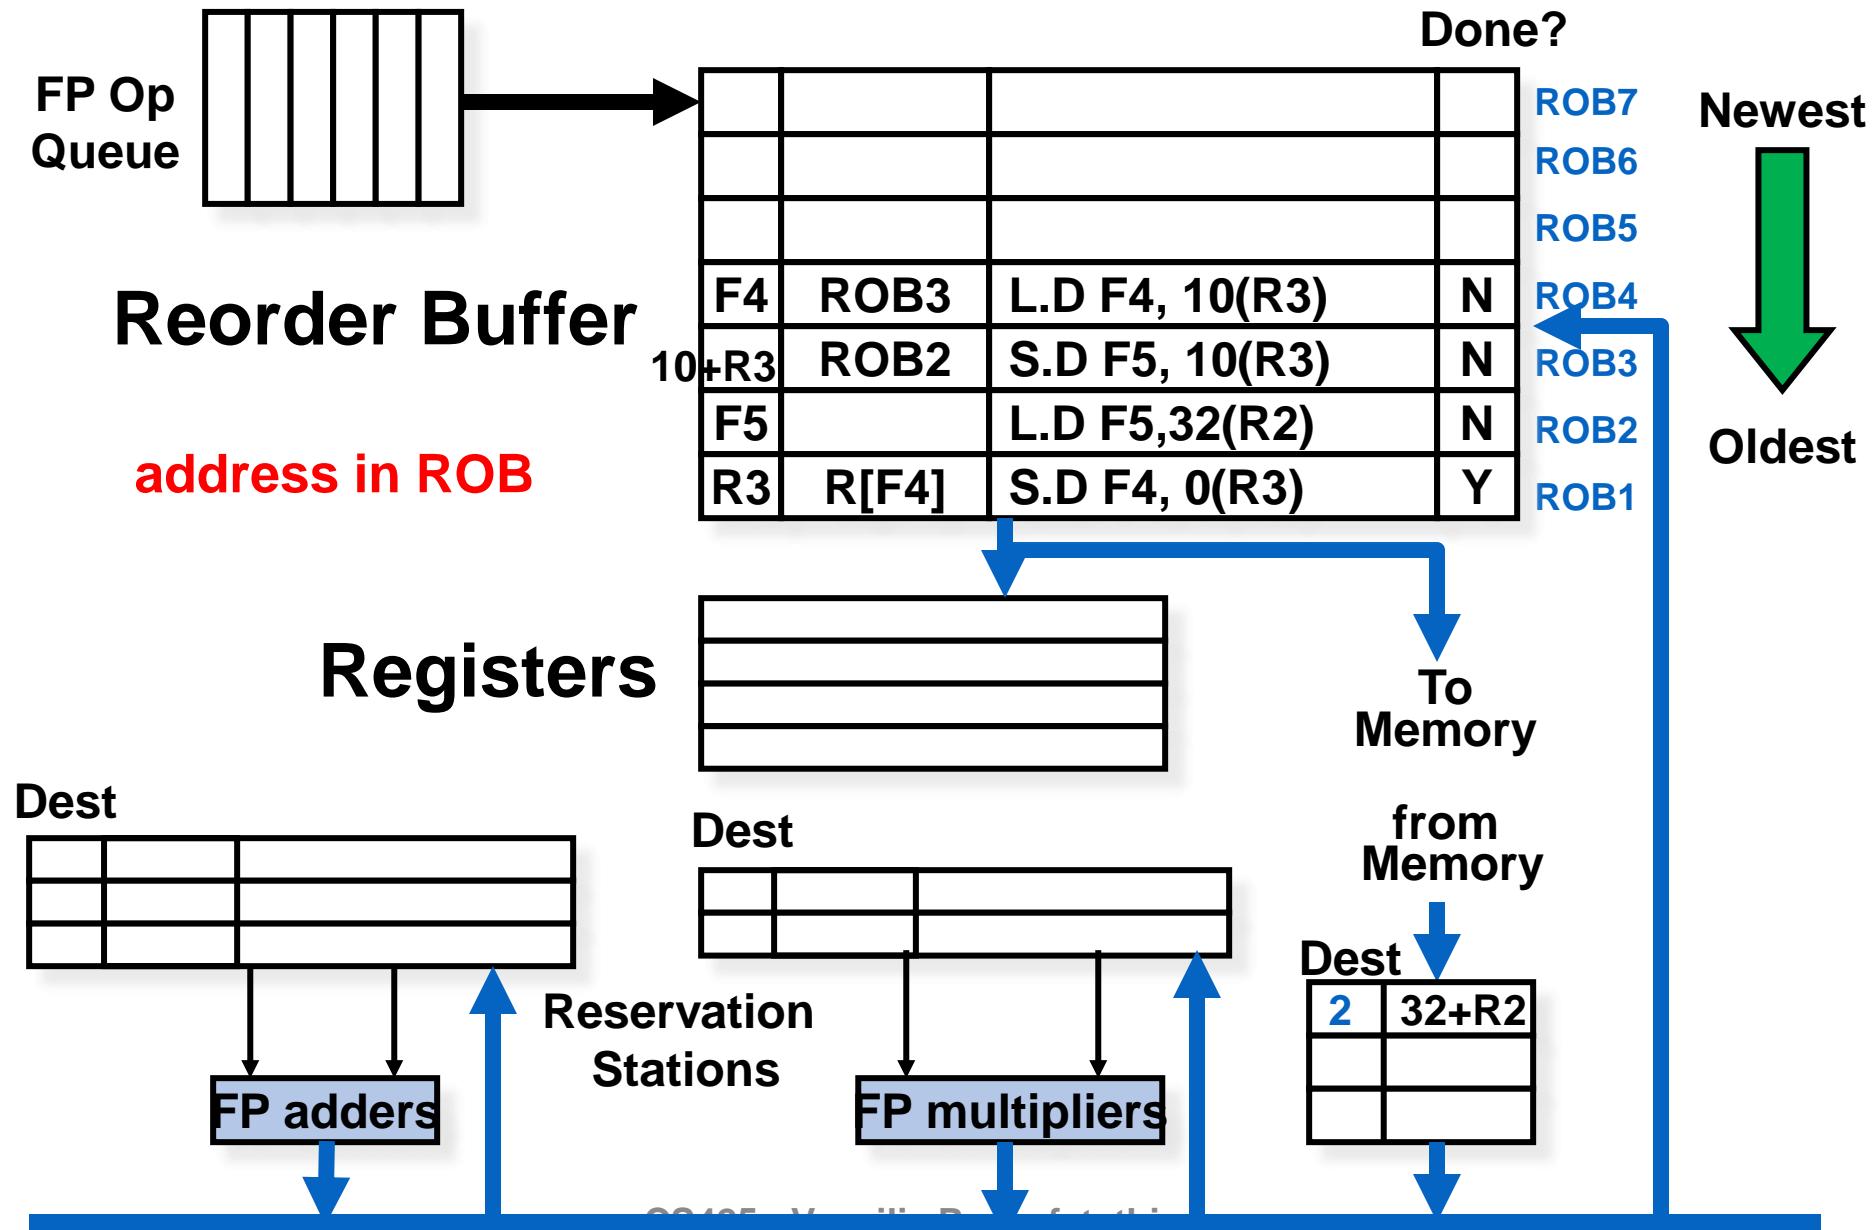
- Challenge: Given a load that follows a store in program order, are these two related?
 - What if there is a RAW hazard between the store and the load?
Eg:

SD	R5, 0 (R2)
LD	R6, 0 (R3)
- Can we proceed and issue the load to the memory system?
 - Store address could be delayed for a long time by some calculation that leads to R2 (e.g. divide).
 - We might want to issue/begin execution of both operations in same cycle.
 - **Solution1:** Answer is that we are not allowed to start load until we know that address $0(R2) \neq 0(R3)$
 - **Solution2:** We might guess/predict whether or not they are dependent (called “dependence speculation”) and use reorder buffer to fixup if we are wrong.

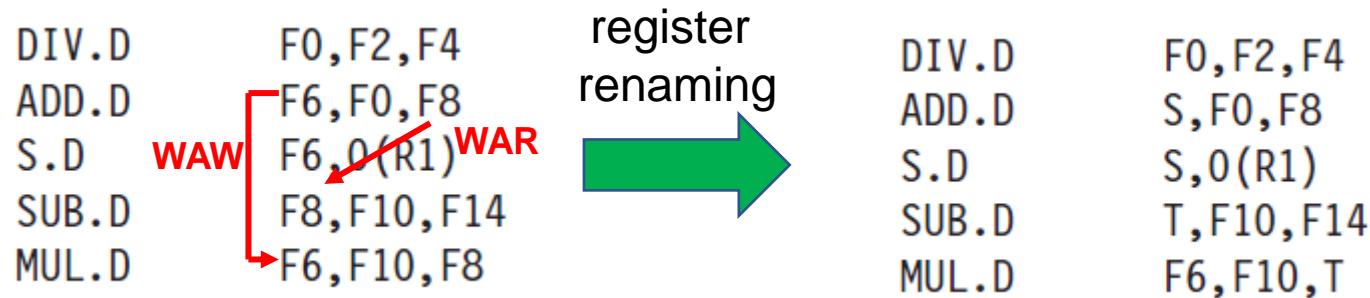
HW support for Memory Disambiguation

- Store buffer keeps all pending stores to memory, in program order
 - Keep track of address (when becomes available) and value (when becomes available)
 - FIFO ordering: will retire stores from this buffer in program order
- When issuing a load, record the head of the store buffer (which stores precede)
- When we have the address of the load, check the buffer:
 - If *any* store prior to load is waiting for its address, stall load
 - If load address matches earlier store address ([associative lookup](#)), then we have a *memory-induced RAW hazard*:
 - store value available \Rightarrow return value
 - store value not available \Rightarrow return ROB number of source
 - Otherwise, send out request to memory
- Stores commit in order, there are no WAW/WAR hazards in memory.

Memory Disambiguation

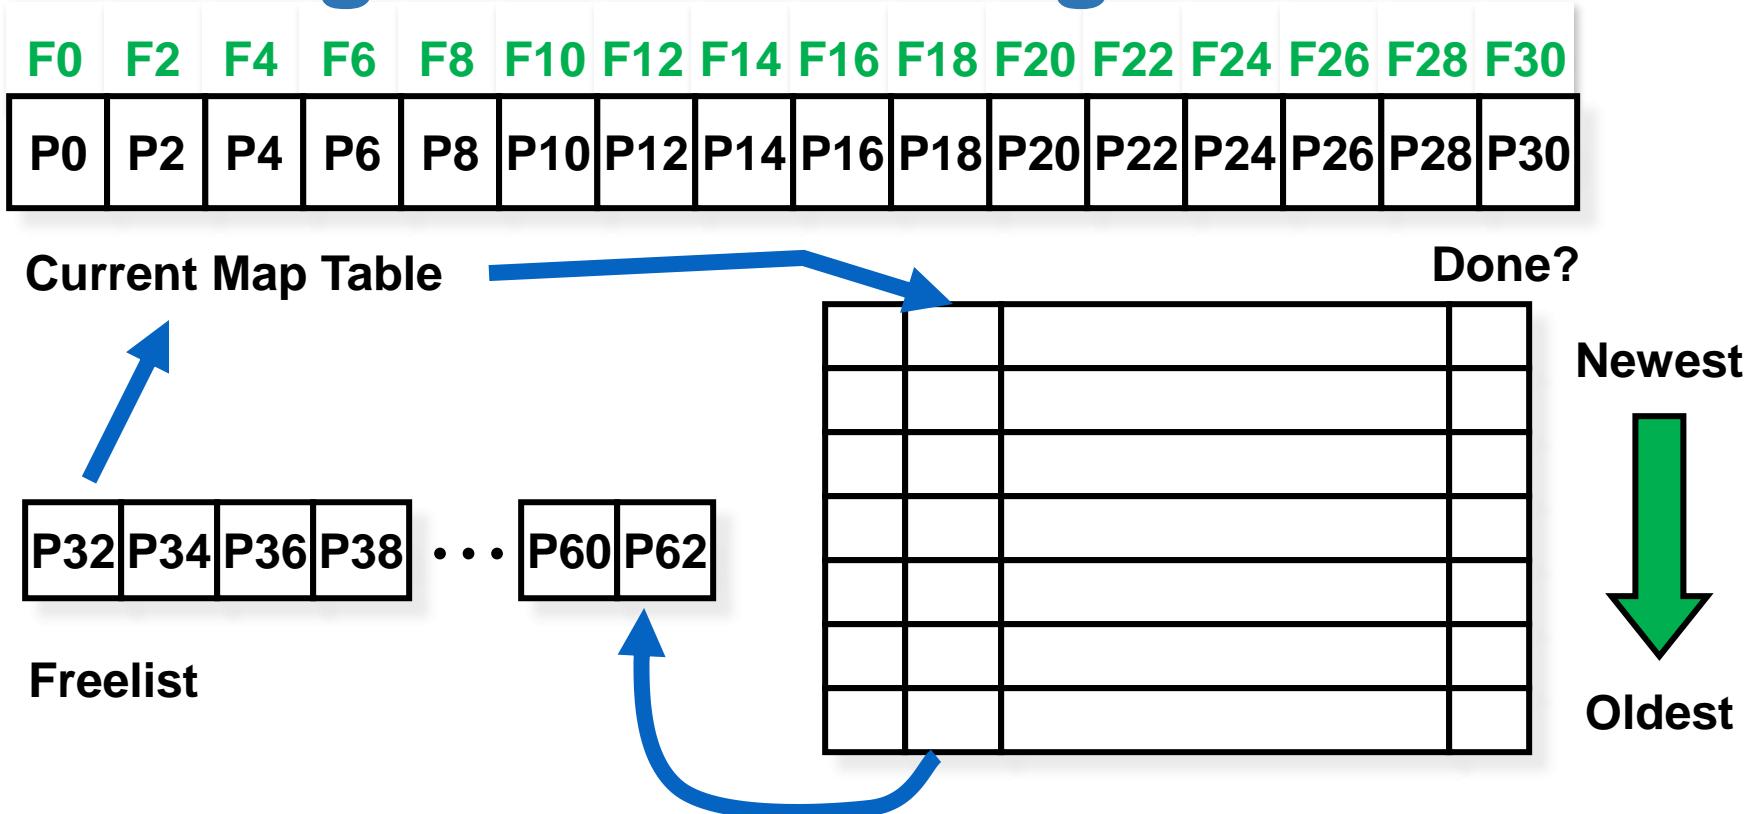


Register Renaming



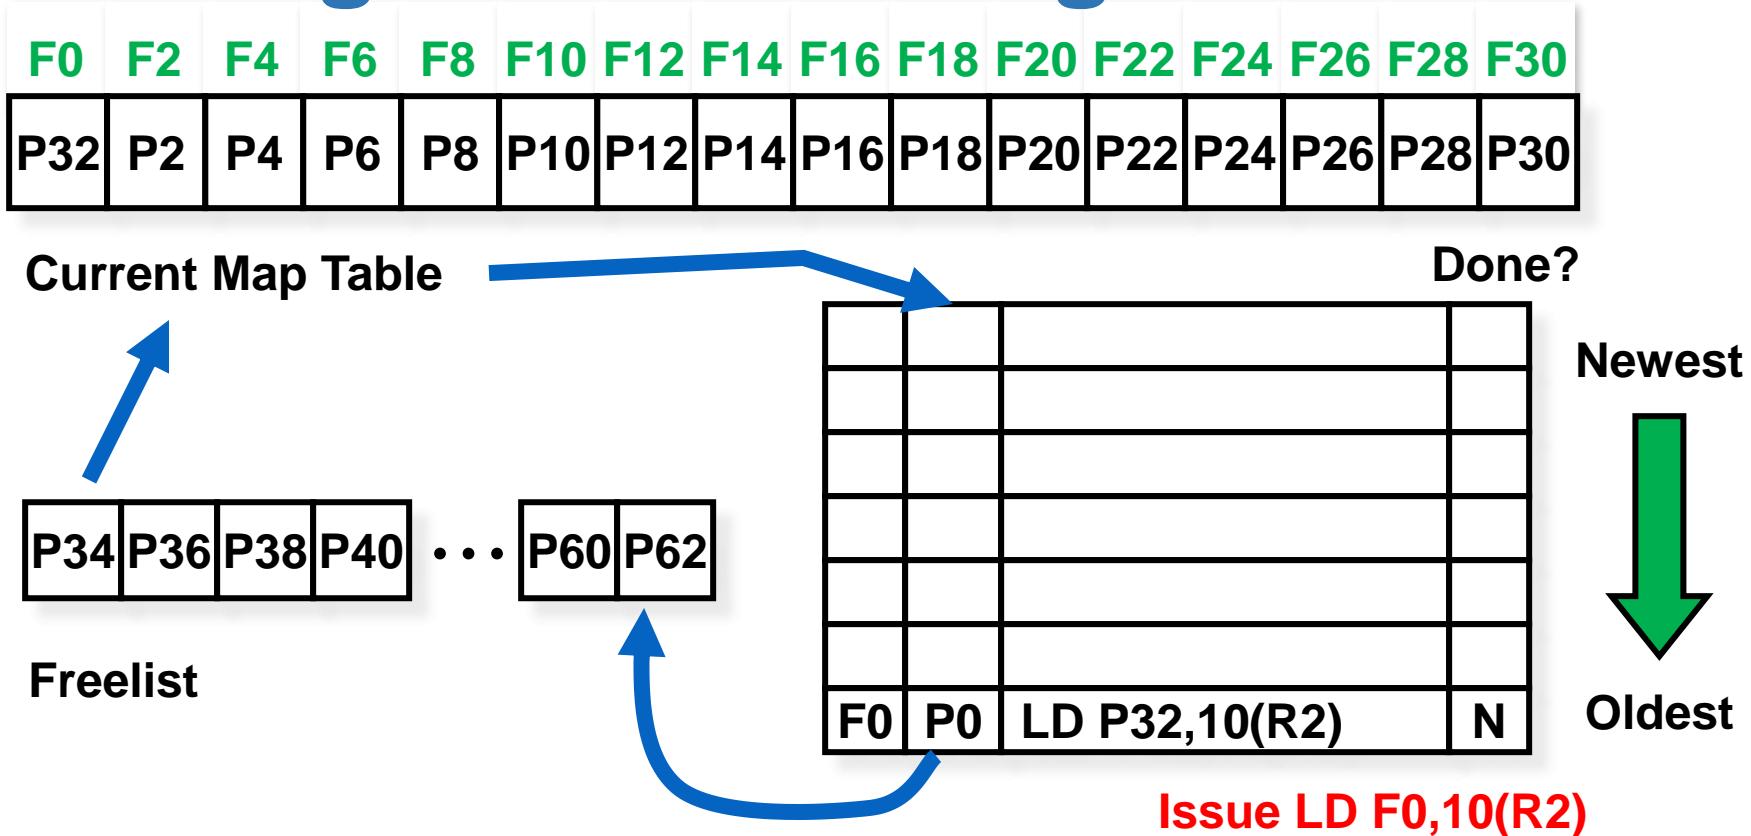
- What happens with branches?
- Tomasulo can handle renaming across branches

Explicit register renaming



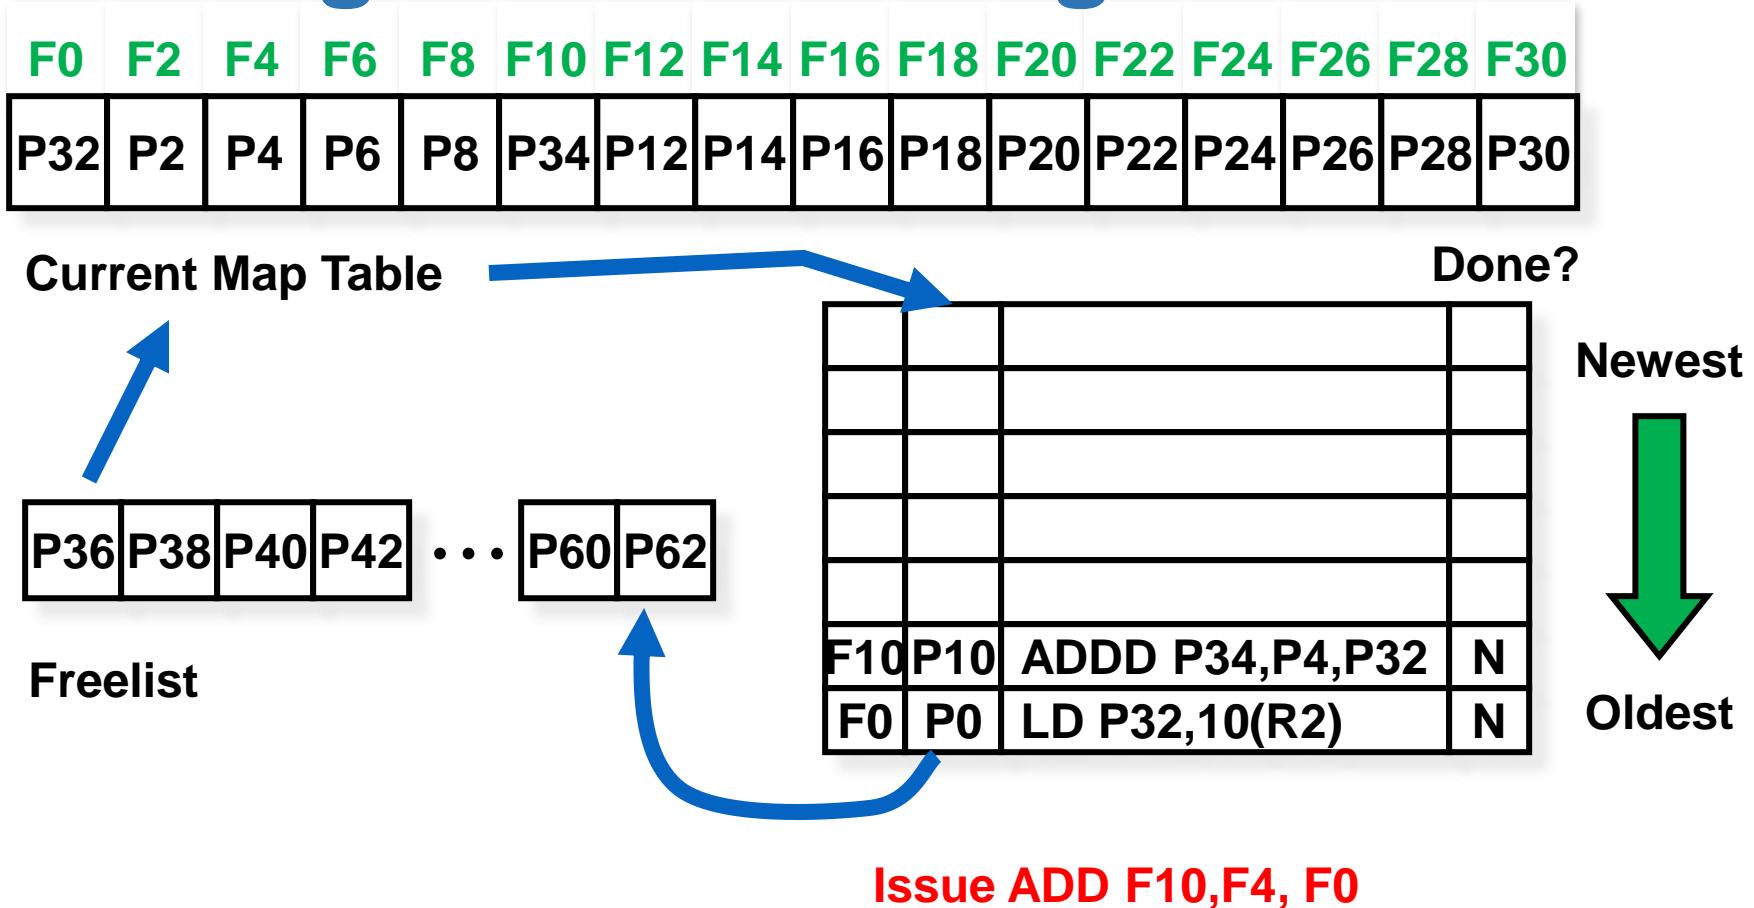
- Hardware equivalent of static, single-assignment (SSA) compiler form
- Physical register file bigger than ISA register file (e.g. 32 Phys regs kai 16 ISA regs)
- Upon issue, every instruction that write a register allocates a new physical register from the freelist

Explicit register renaming

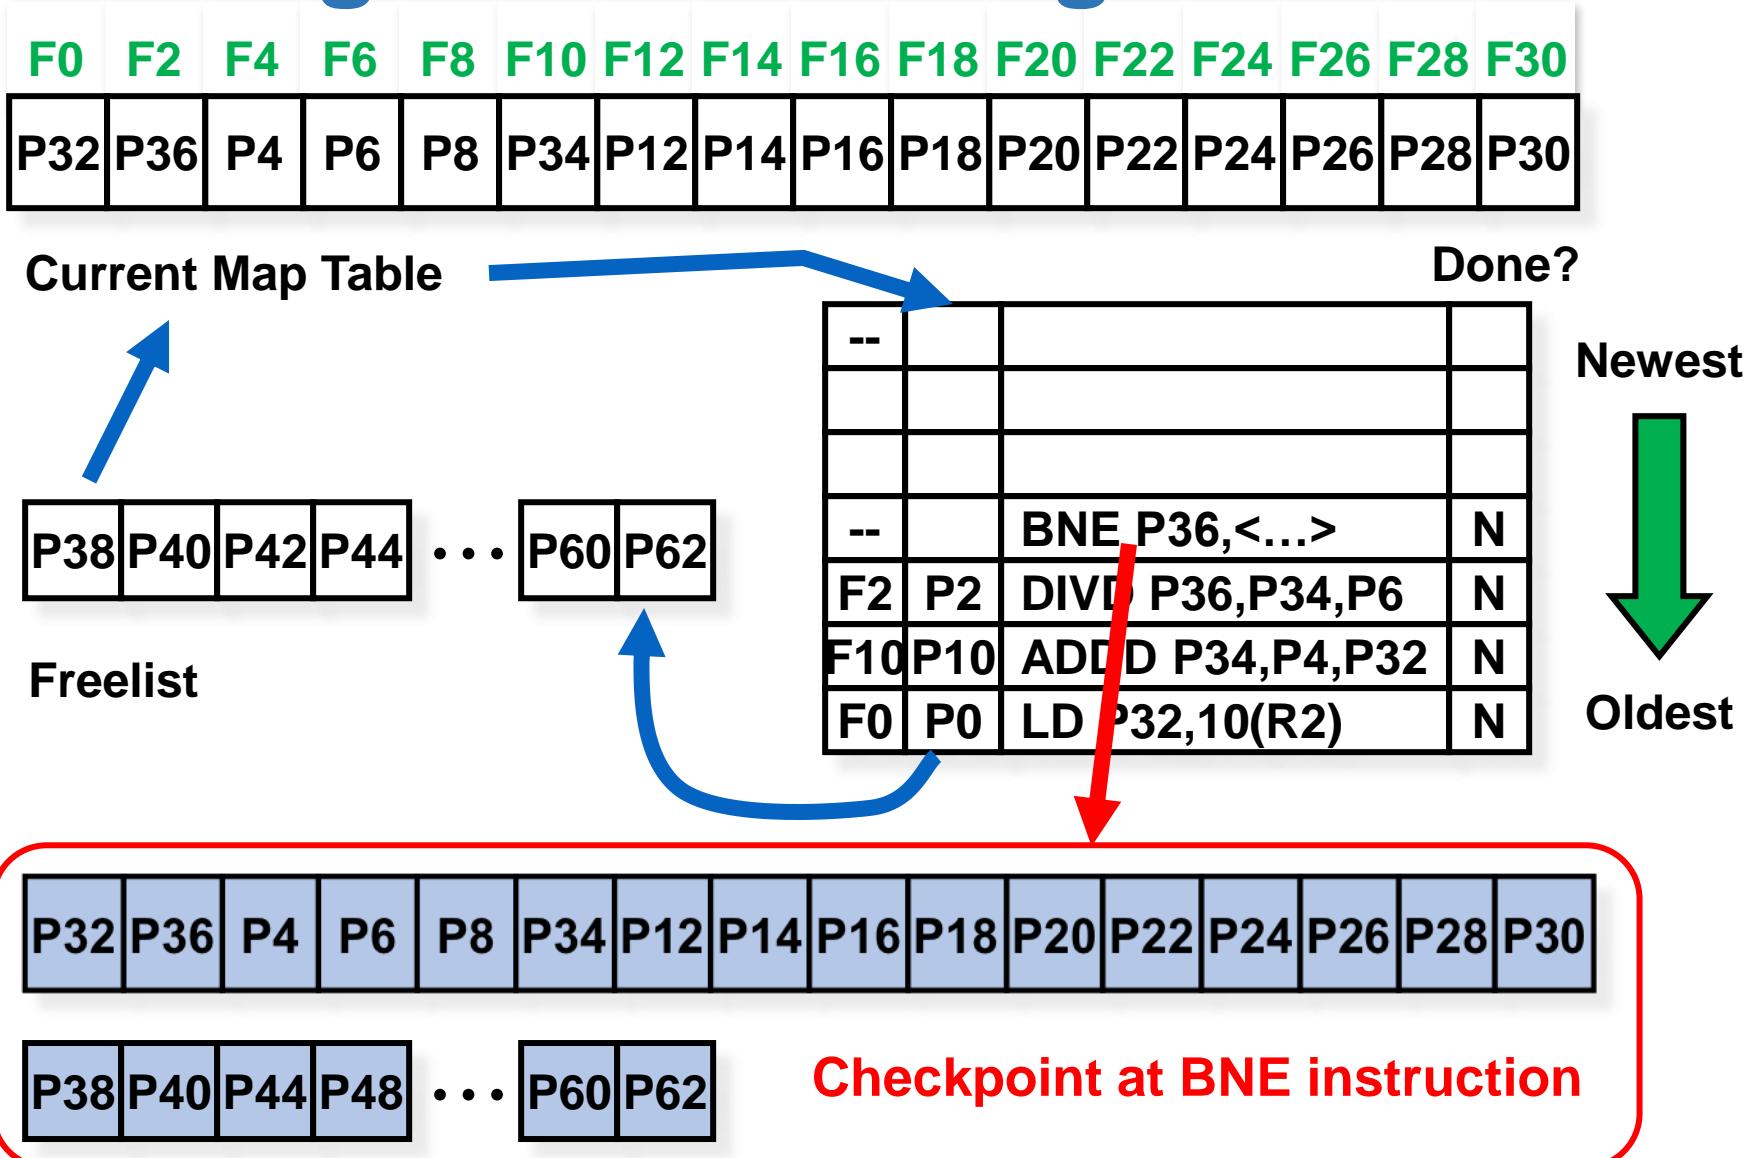


- Note that physical register P0 is “dead” (or not “live”) past the point of this load.
 - When we commit the load, we free up

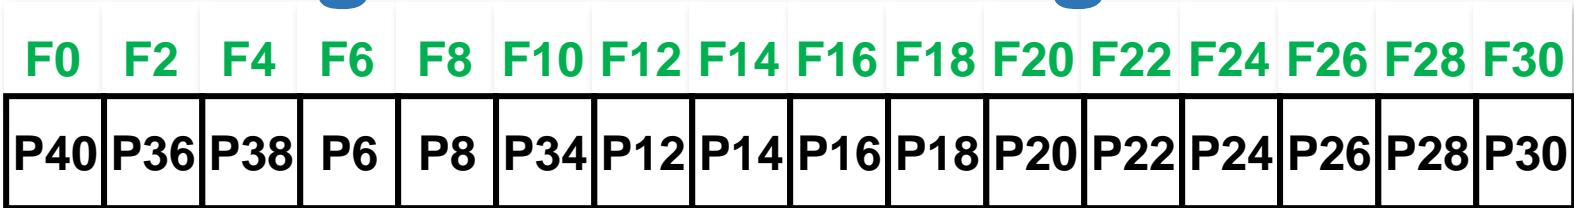
Explicit register renaming



Explicit register renaming



Explicit register renaming



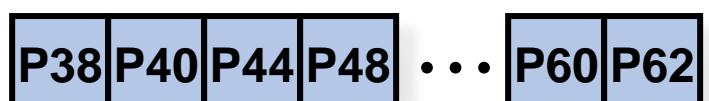
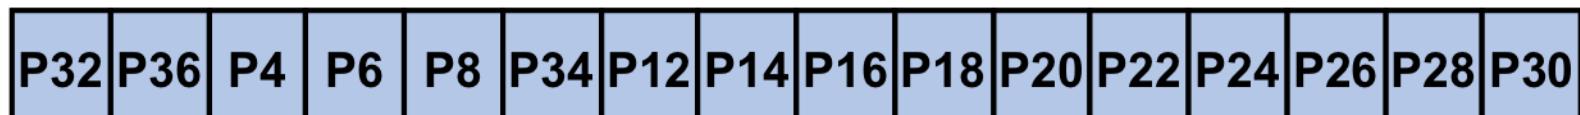
Current Map Table



Freelist

--		SD P40,0(R3)	Y
F0	P32	ADDD P40,P38,P6	Y
F4	P4	LD P38,0(R3)	Y
--		BNE P36,<...>	N
F2	P2	DIVD P36,P34,P6	N
F10	P10	ADDD P34,P4,P32	Y
F0	P0	LD P32,10(R2)	Y

Newest
Oldest



Checkpoint at BNE instruction

Explicit register renaming

